

Fixed-Byte Source-Private Candidate Packets for Bounded Model Collaboration

Anonymous authors

Paper under double-blind review

Abstract

Multi-model language systems usually communicate by emitting text or by sharing internal state such as KV caches. Text is portable but slow and lossy; cache sharing is information-rich but exposes thousands of bytes of model state per token and requires architecture-specific machinery. We study a narrower question: can a frozen source model send a fixed-byte, source-private candidate packet that improves a frozen receiver on public multiple-choice reasoning tasks, without sending source text, hidden states, KV cache, source logits, or gold labels? We present LatentWire, a public-coordinate packet protocol in which source and receiver agree on an anchor/spectral basis over candidate answers, the source emits a tiny signed packet, and the receiver decodes the packet against its own public side information. On ARC-Challenge test, an 8-byte packet reaches 0.344 accuracy across 10 seeds, above target-only 0.265 and same-budget structured text 0.300; on OpenBookQA test, a 3-byte packet reaches 0.378 across 5 seeds, above target-only 0.276 and same-budget text 0.350. The decisive audit is also negative: explicit source-index communication reaches 0.346 on ARC-Challenge and 0.378 on OpenBookQA, with packet-minus-source-index paired lower bounds of -0.008 and -0.006. The positive is therefore compact source-candidate transfer, not evidence synthesis beyond the source or a method that beats source-index communication. We also report that a Phi-3 cross-family source fails the same gate, and learned cache/score connectors on available TinyLlama/Qwen cached artifacts do not repair that cached-artifact gap. The contribution is a reproducible fixed-byte source-private transfer result, a falsification protocol for cross-family claims, and a systems accounting boundary showing that task packets are 69.8x smaller than a one-token 1-bit-per-KV-element accounting floor.

1 Introduction

Future language-model systems will often use more than one model: a cheap model may screen candidates, a specialist may inspect hard cases, and a stronger receiver may combine evidence. The simplest interface is text. Recent cache-level systems instead communicate through model internals, such as fused or selectively shared KV caches (Fu et al., 2026; Shi et al., 2026). These approaches are powerful, but they couple communication to large hidden-state objects and serving substrates.

This paper asks whether a much smaller interface can carry useful source-side candidate evidence. In our setting, both models see the same public question and answer choices. The source performs private computation over those choices. The only communicated object is a short byte packet; the receiver never sees source text, source logits, source hidden vectors, KV cache, or the answer key. This is a side-information problem in the spirit of distributed source coding (Slepian & Wolf, 1973; Wyner & Ziv, 1976): the receiver already has the public prompt and its own computation, so the packet should carry only the source evidence that is useful at the decoder.

Our strongest current result is deliberately scoped. On ARC-Challenge and OpenBookQA (Clark et al., 2018; Mihaylov et al., 2018), same-family Qwen2.5 packet transfer (Yang et al., 2025) improves over target-only and same-budget text controls. However, a




Figure 1: LatentWire protocol. Both sides receive the public task. The source privately scores candidates and emits only a fixed-byte packet. The receiver decodes the packet with a public basis and its own side information. Destructive controls scramble the shared coordinate system or candidate relation.

strict Phi-3 source-family replacement (Abdin et al., 2024) fails, and additional connector repairs fail on the available TinyLlama/Qwen cached artifacts. This boundary matters: the paper is not claiming solved cross-family latent communication.

The contributions are:

- A fixed-byte, source-private packet protocol for multiple-choice candidate transfer with no source text, hidden vector, KV cache, source-logit, or gold-label exposure.
- A public-coordinate anchor/spectral implementation that is positive on ARC-Challenge and OpenBookQA under seed repeats, paired uncertainty, same-budget text controls, and destructive coordinate mismatch controls.
- An explicit source-index/source-rank audit showing that the current positive is mostly source-choice preserving and should be framed as source-candidate transfer, not compression beyond a selected-candidate code.
- A falsification audit showing that the current positive is same-family/source-cache specific: Phi-3 cross-family transfer and cached learned repair connectors fail.
- A systems accounting artifact comparing packet bytes and exposure classes against KV/cache communication and KV quantization floors, while explicitly separating byte accounting from native serving-speed claims.

2 Source-private packet transfer

Let a public multiple-choice instance be $x = (q, c_1, \dots, c_m)$ with answer choices c_i . A source model S and receiver R both observe x , but only S computes the source-side private evidence $e_S(x)$. A packet encoder emits

$$z = E(e_S(x), x), \quad |z| \leq B \text{ bytes.}$$

The receiver predicts

$$\hat{y} = D_R(x, z),$$

where D_R may use public task features and the receiver’s own state, but may not access source text, source hidden states, source KV cache, source logits, or answer labels.

We evaluate a packet only if it beats three surfaces: target-only $D_R(x, \emptyset)$, a same-budget structured text control, and destructive controls that preserve superficial packet size but destroy the intended source-receiver relation. We also audit explicit source-index, source-label-text, source-rank-code, and entropy-matched random-index baselines. We report seed

Benchmark	split	bytes	seeds	src	packet	target	text	pkt-src	pkt-text
ARC-Challenge	test, 1172	8/11	10	0.346	0.344	0.265	0.300	-0.008	+0.025
OpenBookQA	test, 500	3/6	5	0.378	0.378	0.276	0.350	-0.006	+0.000

Table 1: Main audit rows. “Bytes” reports payload/framed bytes. “Src” is explicit answer-key-forbidden source-index accuracy before packet decoding. “Text” is the same-budget structured text control. “Pkt-src” and “pkt-text” are minimum paired-bootstrap 95% lower bounds for packet minus source-index and packet minus same-budget text. Packet-minus-target lower bounds are +0.038 on ARC-Challenge and +0.046 on OpenBookQA.

repeats and paired bootstrap lower bounds. This matters because a packet can be useful while mostly preserving the source’s selected candidate; such a result is source-private transfer, not proof of a shared latent language, reasoning beyond the source endpoint, or compression beyond a one-byte selected-candidate channel.

3 Public-coordinate packets

The current ARC-Challenge implementation uses a public train-anchor chart. For each candidate, we compute a public relative-coordinate feature against a frozen set of training anchors, then project that feature into a low-frequency orthonormal DCT-II spectral basis. The packet is a sparse signed syndrome in this shared coordinate system. The receiver scores candidate compatibility with the packet in the same public basis, combining the packet evidence with receiver-side candidate information.

The design is related to relative representations (Moschella et al., 2023): anchors provide a coordinate system less tied to raw hidden dimensions. Our use is narrower. We do not claim zero-shot latent stitching or global representation alignment; we use a public basis to define a tiny task packet. Random shared anchors work nearly as well as semantic anchors in the current hashed ARC-Challenge implementation, so the correct interpretation is shared coordinate agreement, not semantic anchor names.

OpenBookQA uses the same packet-evaluation contract at a smaller 3-byte budget. Across both tasks, the source-choice cache is answer-key-forbidden: the source sees public questions and choices, not labels. The packet is evaluated on frozen test slices after validation selection.

4 Main results

Table 1 and Figure 2 summarize the current positive evidence and the source-index audit. The ARC-Challenge row uses 10 packet/control seeds and 2000 paired bootstrap samples per seed on 1172 test examples. The OpenBookQA row uses 5 seeds on the 500-example test slice. The source-index audit is generated by reconstructing the same packet rows and adding explicit source-index, source-label-text, source-rank-code, and entropy-matched random-index baselines.

On ARC-Challenge, the matched Fourier/anchor-syndrome packet reaches 0.344 mean test accuracy with minimum seed accuracy 0.342. The answer-key-forbidden source index itself reaches 0.346, and packet predictions match the source-selected candidate at 0.995 mean rate across seeds. Anchor-ID shuffle, anchor-value shuffle, and spectral-bin permutation controls all fail in 0/10 seeds and collapse near target-only. The random shared-anchor diagnostic passes, which weakens any claim that named anchors are semantically special, but strengthens the claim that a shared public coordinate chart is the relevant mechanism.

On OpenBookQA, the 3-byte packet reaches 0.378–0.380 across seeds. The explicit source-index baseline reaches 0.378, and the packet follows the source index at 0.999 mean rate. The largest candidate-derangement control accuracy is 0.212, and the packet beats same-budget structured text in mean by +0.028, with a paired lower bound of +0.000. This second task reduces the risk that the result is only an ARC-Challenge artifact, but it confirms the same




Figure 2: Packet accuracy against target-only and same-budget structured text controls. Error bars show the lower paired-bootstrap lift above target-only; upper uncertainty is omitted for readability because the gate is one-sided.




Figure 3: Payload-rate curve from the source-index audit. Packet payloads below 2 bytes are unsupported by the current sparse signed encoder; the 1-byte source-index baseline is reported separately in Table 1.

112 boundary: the current method is source-choice preserving and same-family, not cross-family
 113 latent reasoning.

114 5 Falsification and failure decomposition

115 The most important negative result is the Phi-3 source-family replacement. We keep the
 116 same ARC-Challenge packet budget, public basis, seed count, and bootstrap budget, but
 117 replace the Qwen source-choice cache with a Phi-3 source. The packet no longer passes. On
 118 the full ARC-Challenge test slice, the Phi-3 packet reaches 0.244 versus target-only 0.265. On
 119 the stricter Qwen-disagreement slice, where the alternate source and Qwen source disagree,
 120 the Phi-3 packet reaches 0.200 versus 0.340 for the Qwen-substituted packet.

Source	surface	n	packet	target	text	ref. pkt	blocker
Phi-3	full test	1172	0.244	0.265	0.232	–	source quality
Phi-3	Qwen-disagree	833	0.200	0.273	0.203	0.340	source quality
TinyLlama	Qwen-disagree	473	0.269	0.268	0.258	0.317	family mismatch
Qwen2.5-1.5B	full test	1172	0.442	0.265	0.401	–	val incomplete

Table 2: Cross-family and source-family diagnostics. “Ref. pkt” is the Qwen-substituted packet on disagreement rows. Qwen2.5-1.5B is encouraging but same-family and validation-incomplete, so it is not a headline cross-family claim.

Method	raw	framed	text	KV/hidden	claim status
LatentWire OpenBookQA packet	3B	6B	no	no	local positive
LatentWire ARC-Challenge packet	8B	11B	no	no	local positive
Same-byte text control	8B	11B	yes	no	control only
1-bit/KV-element floor	768B	768B	no	KV	floor only
KIVI 2-bit KV floor	1536B	1536B	no	KV	floor only
C2C one-token KV floor	12288B	12288B	no	KV	native pending

Table 3: Exposure classes. The paper-ready systems claim is that the communicated object is much smaller and less exposing than KV/cache-state objects under this accounting. It is not a GPU throughput, TTFT, TPOT, or HBM claim.

121 The failure decomposition separates three causes: source endpoint quality, packet corruption,
 122 and common-feature mismatch. For Phi-3, packet predictions follow the Phi-3 source choice
 123 at 0.997, but the source itself is below the target. For TinyLlama, packet fidelity is also
 124 high, but the source-family mismatch causes the disagreement-slice packet to lose to Qwen-
 125 substituted packets. For Qwen2.5-1.5B, the source is stronger and the packet is useful on
 126 frozen test, but the validation disagreement gate is incomplete. Across these rows, the
 127 packet mostly transports the source-selected candidate through the public coordinate basis.
 128 The source-index audit in Table 1 closes one obvious reviewer question: the current packet
 129 does not beat a direct selected-candidate code. The next research gate is therefore a richer
 130 common-feature connector or calibrated source-score packet that can improve beyond
 131 packet-only source choice.

132 We also tested cached learned repairs on the available local artifacts. A candidate-
 133 syndrome connector that sees packet and score-shape features reaches 0.288 versus 0.317
 134 for Qwen-substituted packets on frozen TinyLlama/Qwen disagreement rows. A one-
 135 hidden-layer MLP over cached TinyLlama hidden/query summaries reaches 0.232, below
 136 Qwen-substituted packets, cached Tiny packets, target-only, and same-budget text. These
 137 negatives are useful: simple cached score geometry and mean hidden/query summaries do
 138 not explain the remaining oracle headroom.

139 6 Systems boundary

140 The systems contribution is byte/exposure accounting, not native serving throughput.
 141 Figure 4 compares framed packet bytes with one-token KV/cache-state accounting floors
 142 from neighboring systems and quantization methods. LatentWire sends 6–11 framed bytes
 143 on the two headline rows. A one-token 1-bit-per-KV-element accounting floor is 768 bytes
 144 in our model-configuration accounting; the largest headline packet is therefore 69.8x smaller
 145 than that floor. C2C-style cache transfer, KVComm-style selective sharing, and serving
 146 substrates such as vLLM/PagedAttention and SGLang operate on KV/cache objects rather
 147 than task-level packets (Fu et al., 2026; Shi et al., 2026; Kwon et al., 2023; Zheng et al.,
 148 2024). KV quantization work reduces cache bytes but still communicates or stores model
 149 state (Zandieh et al., 2024; Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2026).




Figure 4: Boundary-byte accounting. Packet rows are measured task-packet objects. KV/cache rows are lower-bound floors or neighboring serving substrates, not native speed comparisons.

7 Related work

Latent and relative representation communication. Relative representations use similarities to anchors to enable latent-space communication under invariances (Moschella et al., 2023). Sparse autoencoder work motivates interpretable feature coordinates (Huben et al., 2024), while feature-space universality work suggests that some features can be related across models (Lan et al., 2024). SAE Bench provides evaluation machinery for sparse autoencoders (Karvonen et al., 2025), but these works do not by themselves provide a packet protocol or task-level evidence transfer result.

Prompt compression and learned prefixes. Prefix tuning and gist tokens learn continuous or discrete prompt-like interfaces within language models (Li & Liang, 2021; Mu et al., 2023). LatentWire differs because the communicated object is a byte packet decoded by a public receiver, not a learned soft prompt injected into one target. Query-bottleneck connectors such as BLIP-2’s Q-Former show that learned bottlenecks can bridge frozen model families across modalities (Li et al., 2023); our failed cached hidden/query summaries suggest that a true tokenwise query-bottleneck is the right future branch, but it is not closed here.

Cache communication and systems. Cache-to-cache and selective KV-sharing systems communicate richer internal state and are the closest competitors for model-to-model communication (Fu et al., 2026; Shi et al., 2026). vLLM and SGLang improve serving through KV-cache management and reuse (Kwon et al., 2023; Zheng et al., 2024). KV quantizers such as QJL, KIVI, KVQuant, and TurboQuant reduce state cost but remain state-transfer or state-storage techniques (Zandieh et al., 2024; Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2026). Our current systems result is a boundary artifact against these state objects, not a native throughput win.

8 Limitations and next gate

The main limitation is scope. The positive results are same-family/source-cache results on public multiple-choice tasks, and the packet largely preserves the source’s selected candidate. We therefore do not claim to beat explicit source-index communication, source-score quantization, or calibrated source-choice baselines; those are required before a stronger compression claim. Phi-3 cross-family replacement fails, TinyLlama common-feature repair fails, and a true cross-family learned query-bottleneck remains future work. We also do not

claim superiority over C2C, KVComm, QJL, KIVI, TurboQuant, vLLM, or SGLang in native serving metrics. Those rows require NVIDIA hardware and matched implementations.

The next exact gate for an ICLR-strength claim is a frozen-endpoint tokenwise connector with an explicit byte or query budget, tested against C2C/KV-style and structured-text baselines, with source-destroying controls and strict same-family versus cross-family separation. The current strongest claim is narrower: fixed-byte source-private packets can transfer useful same-family evidence on ARC-Challenge and OpenBookQA, and the accompanying falsification ladder clarifies where the method stops working.

9 Conclusion

LatentWire shows that a tiny source-private candidate packet can improve frozen receiver decisions on two public QA benchmarks under strict same-family controls. The same evidence also draws a hard boundary: the current method is mostly source-choice preserving, is not a universal latent language, and does not yet solve cross-family communication. This combination of positive packet transfer, negative falsification, and systems byte/exposure accounting is the paper’s core.

References

- Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone, 2024. URL <https://arxiv.org/abs/2404.14219>.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge, 2018. URL <https://arxiv.org/abs/1803.05457>.
- Tianyu Fu, Zihan Min, Hanling Zhang, Jichao Yan, Guohao Dai, Wanli Ouyang, and Yu Wang. Cache-to-cache: Direct semantic communication between large language models. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=LeatkxrBCi>.
- Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 million context length llm inference with kv cache quantization. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-0040. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/028fcbcf85435d39a40c4d61b42c99a4-Abstract-Conference.html.
- Robert Huben, Hoagy Cunningham, Logan Riggs Smith, Aidan Ewart, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=F76bwRSLeK>.
- Adam Karvonen, Can Rager, Johnny Lin, Curt Tigges, Joseph Isaac Bloom, David Chanin, Yeu-Tong Lau, Eoin Farrell, Callum Stuart Mcdougall, Kola Ayonrinde, Demian Till, Matthew Wearden, Arthur Conmy, Samuel Marks, and Neel Nanda. SAEBench: A comprehensive benchmark for sparse autoencoders in language model interpretability. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 29223–29264. PMLR, 2025. URL <https://proceedings.mlr.press/v267/karvonen25a.html>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626, 2023. doi: 10.1145/3600006.3613165. URL <https://dl.acm.org/doi/10.1145/3600006.3613165>.

- 229 Michael Lan, Philip Torr, Austin Meek, Ashkan Khakzar, David Krueger, and Fazl Barez.
230 Quantifying feature space universality across large language models via sparse autoen-
231 coders, 2024. URL <https://arxiv.org/abs/2410.06981>.
- 232 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. BLIP-2: Bootstrapping language-
233 image pre-training with frozen image encoders and large language models. In *Proceedings*
234 *of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine*
235 *Learning Research*, pp. 19730–19742. PMLR, 2023. URL <https://proceedings.mlr.press/v202/li23q.html>.
- 237 Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation.
238 In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and*
239 *the 11th International Conference on Natural Language Processing (Volume 1: Long Papers)*,
240 pp. 4582–4597. Association for Computational Linguistics, 2021. doi: 10.18653/v1/2021.
241 acl-long.353. URL <https://aclanthology.org/2021.acl-long.353/>.
- 242 Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman,
243 Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for kv cache.
244 In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of
245 *Proceedings of Machine Learning Research*, pp. 32332–32344. PMLR, 2024. URL <https://proceedings.mlr.press/v235/liu24bz.html>.
- 247 Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor
248 conduct electricity? a new dataset for open book question answering. In *Proceedings of*
249 *the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2381–2391.
250 Association for Computational Linguistics, 2018. doi: 10.18653/v1/D18-1260. URL
251 <https://aclanthology.org/D18-1260/>.
- 252 Luca Moschella, Valentino Maiorca, Marco Fumero, Antonio Norelli, Francesco Locatello,
253 and Emanuele Rodolà. Relative representations enable zero-shot latent space com-
254 munication. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=Src-nwieGJ>.
- 256 Jesse Mu, Xiang Li, and Noah Goodman. Learning to compress prompts with
257 gist tokens. In *Advances in Neural Information Processing Systems*, volume 36,
258 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/3d77c6dcc7f143aa2154e7f4d5e22d68-Abstract-Conference.html.
- 260 Xiangyu Shi, Marco Chiesa, Gerald Q. Maguire Jr., and Dejan Kostic. KVComm: Enabling
261 efficient llm communication through selective kv sharing. In *International Conference on*
262 *Learning Representations*, 2026. URL <https://openreview.net/forum?id=F7rUng23nw>.
- 263 David Slepian and Jack K. Wolf. Noiseless coding of correlated information sources. *IEEE*
264 *Transactions on Information Theory*, 19(4):471–480, 1973.
- 265 Aaron D. Wyner and Jacob Ziv. The rate-distortion function for source coding with side
266 information at the decoder. *IEEE Transactions on Information Theory*, 22(1):1–10, 1976.
- 267 An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan
268 Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report, 2025. URL
269 <https://arxiv.org/abs/2412.15115>.
- 270 Amir Zandieh, Majid Daliri, and Insu Han. QJL: 1-bit quantized jl transform for kv cache
271 quantization with zero overhead, 2024. URL <https://arxiv.org/abs/2406.03482>.
- 272 Amir Zandieh, Majid Daliri, Majid Hadian, and Vahab Mirrokni. TurboQuant: Online vector
273 quantization with near-optimal distortion rate. In *International Conference on Learning*
274 *Representations*, 2026. URL <https://openreview.net/forum?id=tO3ASKZlok>.
- 275 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao
276 Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett,
277 and Ying Sheng. SGLang: Efficient execution of structured language model pro-
278 grams. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi:

279 10.52202/079017-2000. URL [https://proceedings.neurips.cc/paper_files/paper/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html)
 280 [2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html).

281 A Artifact manifest

- 282 • ARC-Challenge headline: [colm_final/evidence/results/source_private_arc_challenge_fourier_anchor_syndrome_gate_20260502_budget8_10seed_b2000](#).
- 283
- 284 • Phi-3 falsification: [colm_final/evidence/results/source_private_arc_challenge_source_family_cache_falsification_20260502_phi3_cpu_budget8_10seed_b2000](#).
- 285
- 286 • Cross-family decomposition: [colm_final/evidence/results/source_private_arc_cross_family_failure_decomposition_20260502](#).
- 287
- 288 • Systems boundary: [colm_final/evidence/results/source_private_systems_boundary_figure_table_20260502](#).
- 289
- 290 • OpenBookQA headline: [colm_final/evidence/results/source_private_openbookqa_seed_stability_20260501_qwen05_hashed_test_3b](#).
- 291
- 292 • Source-index audit: [colm_final/evidence/results/source_private_colm_acceptance_baselines_20260502](#).
- 293

294 B Reproducibility summary

295 The artifact bundle records exact shell commands, input hashes, output hashes, and rerun
 296 results. The paper rows are produced from cached, answer-key-forbidden source-choice
 297 artifacts with the following settings:

	Row	script	seeds	key settings
298	ARC-Challenge	Fourier/anchor syndrome gate	10	8B, 2000 bootstraps
	OpenBookQA	seed-stability gate	5	3B, hashed features, 500 bootstraps
	Source-index audit	acceptance baseline audit	10/5	1000 bootstraps, 2/3/4/8B rate curve
	Phi-3	source-family cache falsification	10	8B, cached Phi-3 source
	Systems	systems boundary table	–	byte/exposure accounting only

299 Frozen artifacts are hash-verified. In the final local pass, OpenBookQA, systems, and failure-
 300 decomposition reruns completed; the more expensive full ARC-Challenge/Phi-3 reruns
 301 are represented by frozen hashes unless rerun offline. Rerun JSON files are metric-oriented
 302 rather than byte-identical because they contain creation times and some artifacts include
 303 measured local latencies.

304 C Lay explanation of the main experiment

305 The experiment asks whether one model can send a tiny hint to another model. Both models
 306 see the same question and answer choices. The sender chooses a compressed coordinate hint,
 307 only a few bytes long. The receiver uses that hint to choose an answer. If both sides agree on
 308 the coordinate system, the hint helps. If we scramble the coordinate system, the hint stops
 309 helping. That is evidence for a real shared packet interface on the tested same-family setting,
 310 while the failed Phi-3 run shows the interface is not yet general across model families.