

Fixed-Byte Source-Private Packets for Cross-Model Evidence Transfer

Anonymous authors

Paper under double-blind review

Abstract

Multi-model language systems usually communicate by emitting text or by sharing internal state such as KV caches. Text is portable but slow and lossy; cache sharing is information-rich but exposes thousands of bytes of model state per token and requires architecture-specific machinery. We study a narrower question: can a frozen source model send a fixed-byte, source-private evidence packet that improves a frozen receiver on public multiple-choice reasoning tasks, without sending source text, hidden states, KV cache, or answer labels? We present LatentWire, a public-coordinate packet protocol in which source and receiver agree on an anchor/spectral basis over candidate answers, the source emits a tiny signed packet, and the receiver decodes the packet against its own public side information. On ARC-Challenge test, an 8-byte packet reaches 0.344 accuracy across 10 seeds, above target-only 0.265 and same-byte structured text 0.300, with destructive coordinate controls failing in all 10 seeds. On OpenBookQA test, a 3-byte packet reaches 0.378 across 5 seeds, above target-only 0.276 and same-byte text 0.350. We also report strict negative evidence: a Phi-3 cross-family source fails the same gate, and learned Mac-local cache/score connectors do not repair it. The contribution is therefore not a universal latent language claim. It is a reproducible fixed-byte source-private transfer result, a falsification protocol for cross-family claims, and a systems accounting boundary showing that task packets are 69.8x smaller than even a one-token 1-bit KV-state floor under our current accounting.

1 Introduction

Future language-model systems will often use more than one model: a cheap model may screen candidates, a specialist may inspect hard cases, and a stronger receiver may combine evidence. The simplest interface is text. Recent cache-level systems instead communicate through model internals, such as fused or selectively shared KV caches (Fu et al., 2025; Shi et al., 2025). These approaches are powerful, but they couple communication to large hidden-state objects and serving substrates.

This paper asks whether a much smaller interface can carry useful source-side evidence. In our setting, both models see the same public question and answer choices. The source performs private computation over those choices. The only communicated object is a short byte packet; the receiver never sees source text, source logits, source hidden vectors, KV cache, or the answer key. This is a side-information problem in the spirit of distributed source coding (Slepian & Wolf, 1973; Wyner & Ziv, 1976): the receiver already has the public prompt and its own computation, so the packet should carry only the source evidence that is useful at the decoder.

Our strongest current result is deliberately scoped. On ARC-Challenge and OpenBookQA (Clark et al., 2018; Mihaylov et al., 2018), same-family Qwen2.5 packet transfer (Yang et al., 2025) improves over target-only and same-byte text controls. However, a strict Phi-3 source-family replacement (Abdin et al., 2024) fails, and additional connector repairs fail on the available Mac-local caches. This boundary matters: the paper is not claiming solved cross-family latent communication.




Figure 1: LatentWire protocol. Both sides receive the public task. The source privately scores candidates and emits only a fixed-byte packet. The receiver decodes the packet with a public basis and its own side information. Destructive controls scramble the shared coordinate system or candidate relation.

44 The contributions are:

- 45 • A fixed-byte, source-private packet protocol for multiple-choice evidence transfer
- 46 with no source text, hidden vector, KV cache, or answer-label exposure.
- 47 • A public-coordinate anchor/spectral implementation that is positive on ARC-
- 48 Challenge and OpenBookQA under seed repeats, paired uncertainty, same-byte
- 49 text controls, and destructive coordinate mismatch controls.
- 50 • A falsification ladder showing that the current positive is same-family/source-cache
- 51 specific: Phi-3 cross-family transfer and cached learned repair connectors fail.
- 52 • A systems accounting artifact comparing packet bytes and exposure classes against
- 53 KV/cache communication and KV quantization floors, while explicitly separating
- 54 byte accounting from native serving-speed claims.

55 2 Source-private packet transfer

56 Let a public multiple-choice instance be $x = (q, c_1, \dots, c_m)$ with answer choices c_i . A source

57 model S and receiver R both observe x , but only S computes the source-side private evidence

58 $e_S(x)$. A packet encoder emits

$$z = E(e_S(x), x), \quad |z| \leq B \text{ bytes.}$$

59 The receiver predicts

$$\hat{y} = D_R(x, z),$$

60 where D_R may use public task features and the receiver’s own state, but may not access

61 source text, source hidden states, source KV cache, source logits, or answer labels.

62 We evaluate a packet only if it beats three surfaces: target-only $D_R(x, \emptyset)$, a same-byte

63 structured text control, and destructive controls that preserve superficial packet size but

64 destroy the intended source-receiver relation. We report seed repeats and paired bootstrap

65 lower bounds, and we separate same-family positives from strict cross-family replacement.

66 3 Public-coordinate packets

67 The current ARC-Challenge implementation uses a public train-anchor chart. For each

68 candidate, we compute a public relative-coordinate feature against a frozen set of training

69 anchors, then project that feature into a low-frequency orthonormal DCT-II spectral basis.

Benchmark	split	bytes	seeds	packet	target	text	CI low
ARC-Challenge	test, 1172	8/11	10/10	0.344	0.265	0.300	+0.038
OpenBookQA	test, 500	3/6	5/5	0.378	0.276	0.350	+0.038

Table 1: Main positive rows. “Bytes” reports payload/framed bytes. “Text” is the same-byte structured text control. “CI low” is the minimum paired-bootstrap 95% lower bound for packet minus target-only.




Figure 2: Packet accuracy against target-only and same-byte structured text controls. Error bars show the lower paired-bootstrap lift above target-only; upper uncertainty is omitted for readability because the gate is one-sided.

70 The packet is a sparse signed syndrome in this shared coordinate system. The receiver
71 scores candidate compatibility with the packet in the same public basis, combining the
72 packet evidence with receiver-side candidate information.

73 The design is related to relative representations (Moschella et al., 2023): anchors provide a
74 coordinate system less tied to raw hidden dimensions. Our use is narrower. We do not claim
75 zero-shot latent stitching or global representation alignment; we use a public basis to define
76 a tiny task packet. Random shared anchors work nearly as well as semantic anchors in
77 the current hashed ARC-Challenge implementation, so the correct interpretation is shared
78 coordinate agreement, not semantic anchor names.

79 OpenBookQA uses the same packet-evaluation contract at a smaller 3-byte budget. Across
80 both tasks, the source-choice cache is answer-key-forbidden: the source sees public questions
81 and choices, not labels. The packet is evaluated on frozen test slices after validation selection.

82 4 Main results

83 Table 1 and Figure 2 summarize the current COLM-ready positive evidence. The ARC-
84 Challenge row uses 10 packet/control seeds and 2000 paired bootstrap samples per seed on
85 1172 test examples. The OpenBookQA row uses 5 seeds on the 500-example test slice.

86 On ARC-Challenge, the matched Fourier/anchor-syndrome packet reaches 0.344 mean test
87 accuracy with minimum seed accuracy 0.342. Anchor-ID shuffle, anchor-value shuffle, and
88 spectral-bin permutation controls all fail in 0/10 seeds and collapse near target-only. The
89 random shared-anchor diagnostic passes, which weakens any claim that named anchors are
90 semantically special, but strengthens the claim that a shared public coordinate chart is the
91 relevant mechanism.

Source	surface	n	packet	target	text	ref. pkt	blocker
Phi-3	full test	1172	0.244	0.265	0.232	–	source quality
Phi-3	Qwen-disagree	833	0.200	0.273	0.203	0.340	source quality
TinyLlama	Qwen-disagree	473	0.269	0.268	0.258	0.317	family mismatch
Qwen2.5-1.5B	full test	1172	0.442	0.265	0.401	–	val incomplete

Table 2: Cross-family and source-family diagnostics. “Ref. pkt” is the Qwen-substituted packet on disagreement rows. Qwen2.5-1.5B is encouraging but same-family and validation-incomplete, so it is not a headline cross-family claim.

On OpenBookQA, the 3-byte packet reaches 0.378–0.380 across seeds. The largest candidate-derangement control accuracy is 0.212, and the minimum lift over same-byte structured text is +0.028. This second task matters because it reduces the risk that the result is only an ARC-Challenge artifact, though it is still a same-family packet result rather than a cross-family result.

5 Falsification and failure decomposition

The most important negative result is the Phi-3 source-family replacement. We keep the same ARC-Challenge packet budget, public basis, seed count, and bootstrap budget, but replace the Qwen source-choice cache with a Phi-3 source. The packet no longer passes. On the full ARC-Challenge test slice, the Phi-3 packet reaches 0.244 versus target-only 0.265. On the stricter Qwen-disagreement slice, where the alternate source and Qwen source disagree, the Phi-3 packet reaches 0.200 versus 0.340 for the Qwen-substituted packet.

The failure decomposition separates three causes: source endpoint quality, packet corruption, and common-feature mismatch. For Phi-3, packet predictions follow the Phi-3 source choice at 0.997, but the source itself is below the target. For TinyLlama, packet fidelity is also high, but the source-family mismatch causes the disagreement-slice packet to lose to Qwen-substituted packets. For Qwen2.5-1.5B, the source is stronger and the packet is useful on frozen test, but the validation disagreement gate is incomplete. The next research gate is therefore not another scalar packet shuffle; it is a stronger source or a richer common-feature connector.

We also tested cached learned repairs on the Mac-local artifacts. A candidate-syndrome connector that sees packet and score-shape features reaches 0.288 versus 0.317 for Qwen-substituted packets on frozen TinyLlama/Qwen disagreement rows. A one-hidden-layer MLP over cached TinyLlama hidden/query summaries reaches 0.232, below Qwen-substituted packets, cached Tiny packets, target-only, and same-byte text. These negatives are useful: simple cached score geometry and mean hidden/query summaries do not explain the remaining oracle headroom.

6 Systems boundary

The systems contribution is byte/exposure accounting, not native serving throughput. Figure 3 compares framed packet bytes with one-token KV/cache-state floors from neighboring systems and quantization methods. LatentWire sends 6–11 framed bytes on the two headline rows. A one-token 1-bit KV-state floor is 768 bytes in our accounting; the largest headline packet is therefore 69.8x smaller than that floor. C2C-style cache transfer, KVComm-style selective sharing, and serving substrates such as vLLM/PagedAttention and SGLang operate on KV/cache objects rather than task-level packets (Fu et al., 2025; Shi et al., 2025; Kwon et al., 2023; Zheng et al., 2023). KV quantization work reduces cache bytes but still communicates or stores model state (Zandieh et al., 2024; Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2025).




Figure 3: Boundary-byte accounting. Packet rows are measured task-packet objects. KV/cache rows are lower-bound floors or neighboring serving substrates, not native speed comparisons.

Method	raw	framed	text	KV/hidden	claim status
LatentWire OpenBookQA packet	3B	6B	no	no	local positive
LatentWire ARC-Challenge packet	8B	11B	no	no	local positive
Same-byte text control	8B	11B	yes	no	control only
QJL 1-bit KV floor	768B	768B	no	KV	floor only
KIVI 2-bit KV floor	1536B	1536B	no	KV	floor only
C2C one-token KV floor	12288B	12288B	no	KV	native pending

Table 3: Exposure classes. The paper-ready systems claim is that the communicated object is much smaller and less exposing than KV/cache-state objects under this accounting. It is not a GPU throughput, TTFT, TPOT, or HBM claim.

130 7 Related work

131 **Latent and relative representation communication.** Relative representations use similarities
 132 ties to anchors to enable latent-space communication under invariances (Moschella et al.,
 133 2023). Sparse autoencoders and feature-space universality work motivate common-feature
 134 interfaces, but they do not by themselves provide a packet protocol or task-level evidence
 135 transfer result (Cunningham et al., 2023; Engels et al., 2024; Karvonen et al., 2025).

136 **Prompt compression and learned prefixes.** Prefix tuning and gist tokens learn continuous
 137 or discrete prompt-like interfaces within language models (Li & Liang, 2021; Mu et al., 2023).
 138 LatentWire differs because the communicated object is a byte packet decoded by a public
 139 receiver, not a learned soft prompt injected into one target. Query-bottleneck connectors
 140 such as BLIP-2’s Q-Former show that learned bottlenecks can bridge frozen model families
 141 across modalities (Li et al., 2023); our failed Mac-local hidden/query summaries suggest
 142 that a true tokenwise query-bottleneck is the right future branch, but it is not closed here.

143 **Cache communication and systems.** Cache-to-cache and selective KV-sharing systems
 144 communicate richer internal state and are the closest competitors for model-to-model
 145 communication (Fu et al., 2025; Shi et al., 2025). vLLM and SGLang improve serving
 146 through KV-cache management and reuse (Kwon et al., 2023; Zheng et al., 2023). KV
 147 quantizers such as QJL, KIVI, KVQuant, and TurboQuant reduce state cost but remain
 148 state-transfer or state-storage techniques (Zandieh et al., 2024; Liu et al., 2024; Hooper et al.,

2024; Zandieh et al., 2025). Our current systems result is a boundary artifact against these state objects, not a native throughput win.

8 Limitations and next gate

The main limitation is scope. The positive results are same-family/source-cache results on public multiple-choice tasks. Phi-3 cross-family replacement fails, TinyLlama common-feature repair fails, and a true cross-family learned query-bottleneck remains future work. We also do not claim superiority over C2C, KVComm, QJL, KIVI, TurboQuant, vLLM, or SGLang in native serving metrics. Those rows require NVIDIA hardware and matched implementations.

The next exact gate for an ICLR-strength claim is a frozen-endpoint tokenwise connector with an explicit byte or query budget, tested against C2C/KV-style and structured-text baselines, with source-destroying controls and strict same-family versus cross-family separation. For COLM, the current strongest submission is the narrower claim: fixed-byte source-private packets can transfer useful same-family evidence on ARC-Challenge and OpenBookQA, and the accompanying falsification ladder clarifies where the method stops working.

9 Conclusion

LatentWire shows that a tiny source-private packet can improve frozen receiver decisions on two public QA benchmarks under strict same-family controls. The same evidence also draws a hard boundary: the current method is not a universal latent language and does not yet solve cross-family communication. This combination of positive packet transfer, negative falsification, and systems byte/exposure accounting is the paper’s current COLM-ready core.

References

- Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone, 2024. URL <https://arxiv.org/abs/2404.14219>.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge, 2018. URL <https://arxiv.org/abs/1803.05457>.
- Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models, 2023. URL <https://arxiv.org/abs/2309.08600>.
- Joshua Engels, Eric J. Michaud, Isaac Liao, Wes Gurnee, and Max Tegmark. Quantifying feature space universality across large language models via sparse autoencoders, 2024. URL <https://arxiv.org/abs/2410.06981>.
- Tianyu Fu, Zihan Min, Hanling Zhang, Jichao Yan, Guohao Dai, Wanli Ouyang, and Yu Wang. Cache-to-cache: Direct semantic communication between large language models, 2025. URL <https://arxiv.org/abs/2510.03215>.
- Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Hyounjun Genc, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 million context length llm inference with kv cache quantization, 2024. URL <https://arxiv.org/abs/2401.18079>.
- Adam Karvonen, Bilal Chughtai, Lawrence Chan, Tom Lieberum, Samuel Marks, Can Rager, Jeff Wu, and Neel Nanda. SAEBench: A comprehensive benchmark for sparse autoencoders in language model interpretability, 2025. URL <https://arxiv.org/abs/2503.09532>.

- 195 Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu,
196 Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large
197 language model serving with pagedattention. In *ACM Symposium on Operating Systems*
198 *Principles*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- 199 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. BLIP-2: Bootstrapping language-
200 image pre-training with frozen image encoders and large language models. In *International*
201 *Conference on Machine Learning*, 2023. URL <https://arxiv.org/abs/2301.12597>.
- 202 Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation.
203 In *Annual Meeting of the Association for Computational Linguistics*, 2021. URL [https://](https://arxiv.org/abs/2101.00190)
204 arxiv.org/abs/2101.00190.
- 205 Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman,
206 Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for kv cache.
207 In *International Conference on Machine Learning*, 2024. URL [https://arxiv.org/abs/2402.](https://arxiv.org/abs/2402.02750)
208 [02750](https://arxiv.org/abs/2402.02750).
- 209 Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor
210 conduct electricity? a new dataset for open book question answering. In *Conference on*
211 *Empirical Methods in Natural Language Processing*, 2018. URL [https://arxiv.org/abs/](https://arxiv.org/abs/1809.02789)
212 [1809.02789](https://arxiv.org/abs/1809.02789).
- 213 Luca Moschella, Valentino Maiorca, Marco Fumero, Antonio Norelli, Francesco Locatello,
214 and Emanuele Rodolà. Relative representations enable zero-shot latent space com-
215 munication. In *International Conference on Learning Representations*, 2023. URL [https:](https://openreview.net/forum?id=Src-nwieGJ)
216 [/openreview.net/forum?id=Src-nwieGJ](https://openreview.net/forum?id=Src-nwieGJ).
- 217 Jesse Mu, Xiang Lisa Li, and Noah Goodman. Learning to compress prompts with gist
218 tokens, 2023. URL <https://arxiv.org/abs/2304.08467>.
- 219 Xiangyu Shi, Marco Chiesa, Gerald Q. Maguire Jr., and Dejan Kostic. KVComm: Enabling
220 efficient llm communication through selective kv sharing, 2025. URL [https://arxiv.org/](https://arxiv.org/abs/2510.03346)
221 [abs/2510.03346](https://arxiv.org/abs/2510.03346).
- 222 David Slepian and Jack K. Wolf. Noiseless coding of correlated information sources. *IEEE*
223 *Transactions on Information Theory*, 19(4):471–480, 1973.
- 224 Aaron D. Wyner and Jacob Ziv. The rate-distortion function for source coding with side
225 information at the decoder. *IEEE Transactions on Information Theory*, 22(1):1–10, 1976.
- 226 An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan
227 Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report, 2025. URL
228 <https://arxiv.org/abs/2412.15115>.
- 229 Amir Zandieh, Majid Daliri, and Insu Han. QJL: 1-bit quantized jl transform for kv cache
230 quantization with zero overhead, 2024. URL <https://arxiv.org/abs/2406.03482>.
- 231 Amir Zandieh, Majid Daliri, Majid Hadian, and Vahab Mirrokni. TurboQuant: Online vector
232 quantization with near-optimal distortion rate, 2025. URL [https://arxiv.org/abs/2504.](https://arxiv.org/abs/2504.19874)
233 [19874](https://arxiv.org/abs/2504.19874).
- 234 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao Yu,
235 Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying
236 Sheng. SGLang: Efficient execution of structured language model programs, 2023. URL
237 <https://arxiv.org/abs/2312.07104>.

A Artifact manifest

- ARC-Challenge headline: [results / source_private_arc_challenge_fourier_anchor_syndrome_gate_20260502_budget8_10seed_b2000](#).
- Phi-3 falsification: [results/source_private_arc_challenge_source_family_cache_falsification.20260502_phi3_cpu_budget8_10seed_b2000](#).
- Cross-family decomposition: [results/source_private_arc_cross_family_failure_decomposition.20260502](#).
- Systems boundary: [results / source_private_systems_boundary_figure_table_20260502](#).
- OpenBookQA headline: [results / source_private_openbookqa_seed_stability_20260501_qwen05_hashed_test_3b](#).

B Lay explanation of the main experiment

The experiment asks whether one model can send a tiny hint to another model. Both models see the same question and answer choices. The sender chooses a compressed coordinate hint, only a few bytes long. The receiver uses that hint to choose an answer. If both sides agree on the coordinate system, the hint helps. If we scramble the coordinate system, the hint stops helping. That is evidence for a real shared packet interface on the tested same-family setting, while the failed Phi-3 run shows the interface is not yet general across model families.