

LatentWire: Source-Private Candidate-Transfer Packets under Destructive Controls

Anonymous authors

Paper under double-blind review

Abstract

Multi-model language systems usually communicate by emitting text or by sharing internal state such as KV caches. Text is portable but exposes source outputs; cache sharing is information-rich but moves and exposes dense model state. We study a deliberately narrower interface: can a frozen source model send a byte-scale, source-private packet that helps a frozen receiver on public multiple-choice reasoning tasks without sending source text, hidden states, KV cache, source logits, or labels? We present LatentWire, a packet protocol and evaluation framework with explicit byte accounting, a source-private threat model, and destructive controls for source-choice, wrong-row, coordinate-shuffle, same-budget text, and cross-family shortcuts. On ARC-Challenge test, an 8-byte payload / 11-byte framed packet reaches 0.344 accuracy across 10 seeds, above target-only 0.265 and same-budget structured text 0.300; on OpenBookQA test, a 3-byte / 6-byte packet reaches 0.378 across 5 seeds, above target-only 0.276 and same-budget text 0.350. The decisive audit is also limiting: explicit source-index communication reaches 0.346 on ARC-Challenge and 0.378 on OpenBookQA, so the current positive is narrow source-candidate transfer, not evidence synthesis beyond the source or a method that beats selected-candidate communication. We also report that a Phi-3 cross-family source fails and that simple cached connector repairs do not close the gap. The workshop contribution is a reproducible packet-transfer result, a strict falsification protocol, and a systems boundary artifact: headline packets are 6–11 framed bytes, while dense KV/cache-transfer regimes begin hundreds to thousands of bytes per token under our accounting. We do not claim solved latent communication, C2C superiority, or native GPU latency/HBM/energy wins.

1 Introduction

Future language-model systems will often use more than one model: a cheap model may screen candidates, a specialist may inspect hard cases, and a stronger receiver may combine evidence. The simplest interface is text. Recent cache-level systems communicate through model internals, such as fused or selectively shared KV caches (Fu et al., 2026; Shi et al., 2026). These approaches are powerful, but they couple communication to large hidden-state objects and serving substrates.

This paper asks whether a much smaller interface can carry useful source-side candidate evidence. In our setting, both models see the same public question and answer choices. The source performs private computation over those choices. The only communicated object is a short byte packet; the receiver never sees source text, source logits, source hidden vectors, KV cache, or the answer key. This is a side-information problem in the spirit of distributed source coding (Slepian & Wolf, 1973; Wyner & Ziv, 1976): the receiver already has the public prompt and its own computation, so the packet should carry only the source evidence that is useful at the decoder. The scientific contribution is therefore not just a packet row, but a reviewable protocol for asking whether such rows survive the controls that usually create false positives.

Our strongest current result is deliberately scoped. On ARC-Challenge and OpenBookQA (Clark et al., 2018; Mihaylov et al., 2018), same-family Qwen2.5 packet trans-




Figure 1: LatentWire protocol. Both sides receive the public task. The source privately scores candidates and emits only a fixed-byte packet. The receiver decodes the packet with a public basis and its own side information. Destructive controls scramble the shared coordinate system or candidate relation.

fer (Yang et al., 2025) improves over target-only and same-budget text controls. However, a strict Phi-3 source-family replacement (Abdin et al., 2024) fails, and additional connector repairs fail on the available TinyLlama/Qwen cached artifacts. This boundary matters: the paper is not claiming solved cross-family latent communication.

The contributions are:

- A byte-scale, source-private packet protocol and threat model for multiple-choice candidate transfer with no source text, hidden vector, KV cache, source-logit, or gold-label exposure.
- A strict evaluation suite that combines target-only, same-budget text, source-index/rank, coordinate-shuffle, wrong-row, same-source-choice, cross-family, and cached-connector controls.
- Controlled narrow positives on ARC-Challenge and OpenBookQA with seed repeats and paired uncertainty, accompanied by an explicit source-index audit showing that the current effect is mostly source-choice preserving.
- A failure-boundary analysis showing that Phi-3 cross-family transfer and simple cached connector repairs fail, so broad latent-language and cross-family claims are unsupported.
- A systems accounting artifact comparing raw, framed, cacheline-rounded, and dense KV/cache byte regimes while separating measured packet objects from analytical byte floors and future native GPU measurements.

This is useful even with a narrow positive because it separates three questions that are often conflated: whether a compact private packet can help, whether it carries more than the source’s preferred answer, and whether it is cheaper or less exposing than dense state transfer. The current answer is positive for the first question, negative for the second, and accounting-only for the third. That scoped result is the paper’s intended workshop contribution.

Where the current result is useful. The current method is most plausible for same-family cascades, privacy-preserving specialist-to-router hints, and bandwidth-limited candidate reranking where sending source generations, logits, hidden states, or KV cache is undesirable. It is not yet a method for open-ended generation, heterogeneous cross-family collaboration, or reasoning synthesis beyond the source model’s selected candidate.

Control family	What it tests	Current role
target-only / zero-source	target-cache or no-packet effects	required baseline
same-budget visible text	whether text at the same byte budget suffices	required baseline
source-index / rank / score	selected-candidate shortcuts	main claim boundary
wrong-row / row shuffle	row-specific source causality	destructive control
same-source-choice wrong-row	source-choice-only transfer	hardening control
anchor/bin/coordinate shuffle	shared-coordinate dependence	mechanism control
candidate	answer-order artifacts	destructive control
roll/derangement		
source-family substitution	cross-family generality	falsification control
cached connector repair	easy learned receiver rescue	negative diagnostic

Table 1: Control suite. The workshop claim requires packet utility to be interpreted through these controls; the current positive is explicitly bounded by source-index and source-family failures.

2 Source-private packet transfer

Let a public multiple-choice instance be $x = (q, c_1, \dots, c_m)$ with answer choices c_i . A source model S and receiver R both observe x , but only S computes the source-side private evidence $e_S(x)$. A packet encoder emits

$$z = E(e_S(x), x), \quad |z| \leq B \text{ bytes.}$$

The receiver predicts

$$\hat{y} = D_R(x, z),$$

where D_R may use public task features and the receiver’s own state, but may not access source text, source hidden states, source KV cache, source logits, or answer labels.

We evaluate a packet only if it beats three surfaces: target-only $D_R(x, \emptyset)$, a same-budget structured text control, and destructive controls that preserve superficial packet size but destroy the intended source-receiver relation. We also audit explicit source-index, source-label-text, source-rank-code, and entropy-matched random-index baselines. We report seed repeats and paired bootstrap lower bounds. This matters because a packet can be useful while mostly preserving the source’s selected candidate; such a result is source-private transfer, not proof of a shared latent language, reasoning beyond the source endpoint, or compression beyond a one-byte selected-candidate channel.

Threat model. The packet channel is source-private but not source-oblivious. The sender may privately compute over the public prompt and answer choices, then transmit only a bounded byte string. The receiver may use that byte string, the public prompt, and its own computation. The receiver may not observe source generations, chain-of-thought, hidden states, KV cache, logits, score vectors, or labels. A passive packet observer sees only the packet bytes and framing, not source internals or text; however, the current packet can still leak the source’s preferred candidate. We therefore distinguish privacy of source state from secrecy of source choice. A method fails the current workshop claim if its apparent gain is matched by a target-only cache effect, same-budget visible text, explicit source-index/rank/score communication, wrong-row transfer, or answer leakage.

Control suite. Table 1 lists the reviewer-facing controls used to interpret the reported rows. The suite is intentionally stronger than the headline claim: several controls are designed to collapse methods that only copy the source’s preferred answer choice.

3 Public-coordinate packets

The current ARC-Challenge implementation uses a public train-anchor chart. For each candidate, we compute a public relative-coordinate feature against a frozen set of training anchors, then project that feature into a low-frequency orthonormal DCT-II spectral basis. The packet is a sparse signed syndrome in this shared coordinate system. The receiver

Benchmark	split	bytes	seeds	src	packet	target	text	pkt-src	pkt-text
ARC-Challenge	test, 1172	8/11	10	0.346	0.344	0.265	0.300	-0.008	+0.025
OpenBookQA	test, 500	3/6	5	0.378	0.378	0.276	0.350	-0.006	+0.000

Table 2: Main audit rows. “Bytes” reports payload/framed bytes. “Src” is explicit answer-key-forbidden source-index accuracy before packet decoding. “Text” is the same-budget structured text control. “Pkt-src” and “pkt-text” are minimum paired-bootstrap 95% lower bounds for packet minus source-index and packet minus same-budget text. Packet-minus-target lower bounds are +0.038 on ARC-Challenge and +0.046 on OpenBookQA.

scores candidate compatibility with the packet in the same public basis, combining the packet evidence with receiver-side candidate information.

The design is related to relative representations (Moschella et al., 2023): anchors provide a coordinate system less tied to raw hidden dimensions. Our use is narrower. We do not claim zero-shot latent stitching or global representation alignment; we use a public basis to define a tiny task packet. Random shared anchors work nearly as well as semantic anchors in the current hashed ARC-Challenge implementation, so the correct interpretation is shared coordinate agreement, not semantic anchor names.

OpenBookQA uses the same packet-evaluation contract at a smaller 3-byte budget. Across both tasks, the source-choice cache is answer-key-forbidden: the source sees public questions and choices, not labels. The packet is evaluated on frozen test slices after validation selection.

4 Main results

Table 2 and Figure 2 summarize the current positive evidence and the source-index audit. The ARC-Challenge row uses 10 packet/control seeds and 2000 paired bootstrap samples per seed on 1172 test examples. The OpenBookQA row uses 5 seeds on the 500-example test slice. The source-index audit is generated by reconstructing the same packet rows and adding explicit source-index, source-label-text, source-rank-code, and entropy-matched random-index baselines.

Reproducibility card. All headline rows are frozen test-slice evaluations. The ARC-Challenge row uses Qwen2.5-family source/target caches, 10 seeds, and 2000 bootstrap samples; the OpenBookQA row uses the same source-private contract with 5 seeds and 500 bootstrap samples; the source-index audit reruns the same frozen prediction surfaces with explicit index/rank/label baselines. Exact input paths, hashes, scripts, and rerun status are listed in the artifact manifest and regenerated by the COLM.v3 review-packet script.

On ARC-Challenge, the matched Fourier/anchor-syndrome packet reaches 0.344 mean test accuracy with minimum seed accuracy 0.342. The answer-key-forbidden source index itself reaches 0.346, and packet predictions match the source-selected candidate at 0.995 mean rate across seeds. Anchor-ID shuffle, anchor-value shuffle, and spectral-bin permutation controls all fail in 0/10 seeds and collapse near target-only. The random shared-anchor diagnostic passes, which weakens any claim that named anchors are semantically special, but strengthens the claim that a shared public coordinate chart is the relevant mechanism.

On OpenBookQA, the 3-byte packet reaches 0.378–0.380 across seeds. The explicit source-index baseline reaches 0.378, and the packet follows the source index at 0.999 mean rate. The largest candidate-derangement control accuracy is 0.212, and the packet beats same-budget structured text in mean by +0.028, with a paired lower bound of +0.000. This second task reduces the risk that the result is only an ARC-Challenge artifact, but it confirms the same boundary: the current method is source-choice preserving and same-family, not cross-family latent reasoning.

The practical readout is an evidence ladder rather than a single win/loss row. The packets beat target-only and same-budget text on the headline rows; coordinate and candidate-destructive controls prevent treating the result as random packet regularization; explicit




Figure 2: Packet accuracy against target-only and same-budget structured text controls. Error bars show the lower paired-bootstrap lift above target-only; upper uncertainty is omitted for readability because the gate is one-sided.

Benchmark	seeds	pkt-target	pkt-text	pkt-source	provenance
ARC-Challenge	10	+0.038	+0.025	-0.008	frozen test, hash-audited
OpenBookQA	5	+0.046	+0.000	-0.006	frozen test, hash-audited

Table 3: Uncertainty summary for the main audit rows. Values are paired-bootstrap 95% lower bounds for packet accuracy minus the named comparator. The zero lower bound against same-budget text on OpenBookQA is intentionally visible: that row is a narrow secondary positive, not a strong text-baseline separation.

source-index/rank controls prevent overstating the mechanism. In framed-byte terms, the target-only lift is $0.079/11 = 0.0072$ accuracy per byte on ARC-Challenge and $0.102/6 = 0.0170$ accuracy per byte on OpenBookQA. Those rates are useful for comparing packet regimes, but they are not evidence that the current packet is better than a direct selected-candidate code.

Claim boundary. The current packet does not improve over explicit source-index communication: source-index reaches 0.346 on ARC-Challenge and 0.378 on OpenBookQA, while the packet reaches 0.344 and 0.378. We therefore claim source-private fixed-byte candidate transfer under destructive controls, not compression beyond a selected-candidate channel.

Breadth audits. We keep additional model and benchmark breadth in the reviewer packet rather than turning every diagnostic into a headline claim. A frozen benchmark-selection gate also records HellaSwag validation-1024 (Zellers et al., 2019) as a positive bounded validation diagnostic, CommonsenseQA (Talmor et al., 2019) as target-positive but same-byte-text saturated, and SciQ (Welbl et al., 2017) as a source-private bridge contract with no positive packet row yet. A separate latest-model emitter audit covers Qwen3.5, Gemma 4, Granite 3.3, and OLMo-2 source-packet behavior on a synthetic hidden-repair benchmark. These rows improve reviewer visibility into breadth and failure modes, but they are not used to strengthen the main ARC-Challenge/OpenBookQA claim.

5 Falsification and failure decomposition

The most important negative result is the Phi-3 source-family replacement. We keep the same ARC-Challenge packet budget, public basis, seed count, and bootstrap budget, but replace the Qwen source-choice cache with a Phi-3 source. The packet no longer passes. On




Figure 3: Payload-rate curve from the source-index audit. Packet payloads below 2 bytes are unsupported by the current sparse signed encoder; the 1-byte source-index baseline is reported separately in Table 2.

Source	surface	n	packet	target	text	ref. pkt	blocker
Phi-3	full test	1172	0.244	0.265	0.232	–	source quality
Phi-3	Qwen-disagree	833	0.200	0.273	0.203	0.340	source quality
TinyLlama	Qwen-disagree	473	0.269	0.268	0.258	0.317	family mismatch
Qwen2.5-1.5B	full test	1172	0.442	0.265	0.401	–	val incomplete

Table 4: Cross-family and source-family diagnostics. “Ref. pkt” is the Qwen-substituted packet on disagreement rows. Qwen2.5-1.5B is encouraging but same-family and validation-incomplete, so it is not a headline cross-family claim.

the full ARC-Challenge test slice, the Phi-3 packet reaches 0.244 versus target-only 0.265. On the stricter Qwen-disagreement slice, where the alternate source and Qwen source disagree, the Phi-3 packet reaches 0.200 versus 0.340 for the Qwen-substituted packet.

The failure decomposition separates three causes: source endpoint quality, packet corruption, and common-feature mismatch. For Phi-3, packet predictions follow the Phi-3 source choice at 0.997, but the source itself is below the target. For TinyLlama, packet fidelity is also high, but the source-family mismatch causes the disagreement-slice packet to lose to Qwen-substituted packets. For Qwen2.5-1.5B, the source is stronger and the packet is useful on frozen test, but the validation disagreement gate is incomplete. Across these rows, the packet mostly transports the source-selected candidate through the public coordinate basis. The source-index audit in Table 2 closes one obvious reviewer question: the current packet does not beat a direct selected-candidate code. The next research gate is therefore a richer common-feature connector or calibrated source-score packet that can improve beyond packet-only source choice.

We also tested cached learned repairs on the available local artifacts. A candidate-syndrome connector that sees packet and score-shape features reaches 0.288 versus 0.317 for Qwen-substituted packets on frozen TinyLlama/Qwen disagreement rows. A one-hidden-layer MLP over cached TinyLlama hidden/query summaries reaches 0.232, below Qwen-substituted packets, cached Tiny packets, target-only, and same-budget text. These negatives are useful: simple cached score geometry and mean hidden/query summaries do not explain the remaining oracle headroom.




Figure 4: Boundary-byte accounting. Packet rows are measured task-packet objects. KV/cache rows are lower-bound floors or neighboring serving substrates, not native speed comparisons.

Method	raw	framed	text	KV/hidden	claim status
LatentWire OpenBookQA packet	3B	6B	no	no	local positive
LatentWire ARC-Challenge packet	8B	11B	no	no	local positive
Same-byte text control	8B	11B	yes	no	control only
1-bit/KV-element floor	768B	768B	no	KV	floor only
KIVI 2-bit KV floor	1536B	1536B	no	KV	floor only
C2C one-token KV floor	12288B	12288B	no	KV	native pending

Table 5: Exposure classes. Packet rows are measured/accounted local communication objects. KV/cache rows are analytical byte floors or native-pending baselines. The paper-ready systems claim is object size and source-exposure accounting, not GPU throughput, TTFT, TPOT, HBM, energy, or C2C superiority.

193 6 Systems boundary

194 The systems contribution is byte/exposure accounting, not native serving throughput.
 195 Figure 4 compares framed packet bytes with one-token KV/cache-state accounting floors
 196 from neighboring systems and quantization methods. LatentWire sends 6–11 framed bytes
 197 on the two headline rows. A one-token 1-bit-per-KV-element accounting floor is 768 bytes
 198 in our model-configuration accounting; the largest headline packet is therefore 69.8x smaller
 199 than that floor. C2C-style cache transfer, KVComm-style selective sharing, and serving
 200 substrates such as vLLM/PagedAttention and SGLang operate on KV/cache objects rather
 201 than task-level packets (Fu et al., 2026; Shi et al., 2026; Kwon et al., 2023; Zheng et al.,
 202 2024). KV quantization work reduces cache bytes but still communicates or stores model
 203 state (Zandieh et al., 2024; Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2026).

204 We label packet-object bytes as measured/accounted local artifacts and dense-state rows as
 205 analytical or literature-derived byte floors. This distinction is important: the current data
 206 support privacy/exposure and byte-budget comparisons, but not claims about GPU latency,
 207 HBM traffic, energy, or serving throughput.

Regime	Example	Object crossing interface	Difference from LatentWire
dense cache fusion	C2C	projected source KV cache	high-bandwidth, source-state exposing
selective KV sharing KV compression	KVComm KV quantizers	selected KV pairs/layers compressed KV/cache state	cache-level, native pending shrinks dense state, not source-private task packets
serving substrate	vLLM/SGLang	KV reuse and memory management	runtime baseline, not communication protocol
prompt/prefix compression	prefix/gist tokens	prompt-like tokens or parameters	not source-private model-to-model evidence
local shortcut base-lines	source-index/rank/score	selected answer or score sketch	must be beaten for stronger claims

Table 6: Related-work and baseline matrix. LatentWire is positioned as low-byte source-private packet transfer with destructive controls, not as a replacement for dense cache fusion or serving-runtime optimization.

7 Related work

Latent and relative representation communication. Relative representations use similarities to anchors to enable latent-space communication under invariances (Moschella et al., 2023). Sparse autoencoder work motivates interpretable feature coordinates (Huben et al., 2024), while feature-space universality work suggests that some features can be related across models (Lan et al., 2024). Crosscoders, transcoders, and activation-transfer methods are natural future sources of interpretable packet atoms or dense communication baselines. SAEbench provides evaluation machinery for sparse autoencoders (Karvonen et al., 2025), but these works do not by themselves provide a byte-counted packet protocol or task-level evidence transfer result under source-choice controls.

Prompt compression and learned prefixes. Prefix tuning and gist tokens learn continuous or discrete prompt-like interfaces within language models (Li & Liang, 2021; Mu et al., 2023); text-compression systems such as LLMLingua and LongLLMLingua are stronger visible-text baselines for future work (Jiang et al., 2023a;b). LatentWire differs because the communicated object is an opaque byte packet decoded by a public receiver, not a learned soft prompt injected into one target or a compressed visible prompt. Query-bottleneck connectors such as BLIP-2’s Q-Former show that learned bottlenecks can bridge frozen model families across modalities (Li et al., 2023); our failed cached hidden/query summaries suggest that a true tokenwise query-bottleneck is the right future branch, but it is not closed here.

Cache communication and systems. Cache-to-cache and selective KV-sharing systems communicate richer internal state and are the closest competitors for model-to-model communication (Fu et al., 2026; Shi et al., 2026). vLLM and SGLang improve serving through KV-cache management and reuse (Kwon et al., 2023; Zheng et al., 2024). KV quantizers such as QJL, KIVI, KVQuant, and TurboQuant reduce state cost but remain state-transfer or state-storage techniques (Zandieh et al., 2024; Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2026). Table 6 summarizes how these baselines differ from LatentWire. Our current systems result is a boundary artifact against these state objects, not a native throughput win.

8 Limitations and next gate

The main limitation is scope. The positive results are same-family/source-cache results on public multiple-choice tasks, and the packet largely preserves the source’s selected candidate. We therefore do not claim to beat explicit source-index communication, source-score quantization, or calibrated source-choice baselines; those are required before a stronger compression claim. Phi-3 cross-family replacement fails, TinyLlama common-feature repair fails, and a true cross-family learned query-bottleneck remains future work. We also do not claim superiority over C2C, KVComm, QJL, KIVI, TurboQuant, vLLM, or SGLang in native serving metrics. Those rows require NVIDIA hardware and matched implementations.

The next exact gate for an ICLR-strength claim is a frozen-endpoint tokenwise connector with an explicit byte or query budget, tested against C2C/KV-style and structured-text baselines, with source-destroying controls and strict same-family versus cross-family separation. The current strongest claim is narrower: fixed-byte source-private packets can transfer useful same-family evidence on ARC-Challenge and OpenBookQA, and the accompanying falsification ladder clarifies where the method stops working.

9 Conclusion

LatentWire shows that a tiny source-private candidate packet can improve frozen receiver decisions on two public QA benchmarks under strict same-family controls. The same evidence also draws a hard boundary: the current method is mostly source-choice preserving, is not a universal latent language, and does not yet solve cross-family communication. This combination of positive packet transfer, negative falsification, and systems byte/exposure accounting is the paper’s core.

References

- Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone, 2024. URL <https://arxiv.org/abs/2404.14219>.
- Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge, 2018. URL <https://arxiv.org/abs/1803.05457>.
- Tianyu Fu, Zihan Min, Hanling Zhang, Jichao Yan, Guohao Dai, Wanli Ouyang, and Yu Wang. Cache-to-cache: Direct semantic communication between large language models. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=LeatkxrBCi>.
- Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 million context length llm inference with kv cache quantization. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-0040. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/028fcbcf85435d39a40c4d61b42c99a4-Abstract-Conference.html.
- Robert Huben, Hoagy Cunningham, Logan Riggs Smith, Aidan Ewart, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=F76bWRSLeK>.
- Huiqiang Jiang, Qianhui Wu, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. LLMingua: Compressing prompts for accelerated inference of large language models, 2023a. URL <https://arxiv.org/abs/2310.05736>.
- Huiqiang Jiang, Qianhui Wu, Xufang Luo, Dongsheng Li, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. LongLLMingua: Accelerating and enhancing LLMs in long context scenarios via prompt compression, 2023b. URL <https://arxiv.org/abs/2310.06839>.
- Adam Karvonen, Can Rager, Johnny Lin, Curt Tigges, Joseph Isaac Bloom, David Chanin, Yeu-Tong Lau, Eoin Farrell, Callum Stuart McDougall, Kola Ayonrinde, Demian Till, Matthew Wearden, Arthur Conmy, Samuel Marks, and Neel Nanda. SAEBench: A comprehensive benchmark for sparse autoencoders in language model interpretability. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 29223–29264. PMLR, 2025. URL <https://proceedings.mlr.press/v267/karvonen25a.html>.

- 294 Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu,
295 Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large
296 language model serving with pagedattention. In *Proceedings of the 29th Symposium on*
297 *Operating Systems Principles*, pp. 611–626, 2023. doi: 10.1145/3600006.3613165. URL
298 <https://dl.acm.org/doi/10.1145/3600006.3613165>.
- 299 Michael Lan, Philip Torr, Austin Meek, Ashkan Khakzar, David Krueger, and Fazl Barez.
300 Quantifying feature space universality across large language models via sparse autoen-
301 coders, 2024. URL <https://arxiv.org/abs/2410.06981>.
- 302 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. BLIP-2: Bootstrapping language-
303 image pre-training with frozen image encoders and large language models. In *Proceedings*
304 *of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine*
305 *Learning Research*, pp. 19730–19742. PMLR, 2023. URL <https://proceedings.mlr.press/v202/li23q.html>.
- 307 Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation.
308 In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and*
309 *the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*,
310 pp. 4582–4597. Association for Computational Linguistics, 2021. doi: 10.18653/v1/2021.
311 acl-long.353. URL <https://aclanthology.org/2021.acl-long.353/>.
- 312 Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman,
313 Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for kv cache.
314 In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of
315 *Proceedings of Machine Learning Research*, pp. 32332–32344. PMLR, 2024. URL <https://proceedings.mlr.press/v235/liu24bz.html>.
- 317 Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor
318 conduct electricity? a new dataset for open book question answering. In *Proceedings of*
319 *the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2381–2391.
320 Association for Computational Linguistics, 2018. doi: 10.18653/v1/D18-1260. URL
321 <https://aclanthology.org/D18-1260/>.
- 322 Luca Moschella, Valentino Maiorca, Marco Fumero, Antonio Norelli, Francesco Locatello,
323 and Emanuele Rodolà. Relative representations enable zero-shot latent space com-
324 munication. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=SrC-nwieGJ>.
- 326 Jesse Mu, Xiang Li, and Noah Goodman. Learning to compress prompts with
327 gist tokens. In *Advances in Neural Information Processing Systems*, volume 36,
328 2023. URL https://proceedings.neurips.cc/paper_files/paper/2023/hash/3d77c6dcc7f143aa2154e7f4d5e22d68-Abstract-Conference.html.
- 330 Xiangyu Shi, Marco Chiesa, Gerald Q. Maguire Jr., and Dejan Kostic. KVComm: Enabling
331 efficient llm communication through selective kv sharing. In *International Conference on*
332 *Learning Representations*, 2026. URL <https://openreview.net/forum?id=F7rUng23nw>.
- 333 David Slepian and Jack K. Wolf. Noiseless coding of correlated information sources. *IEEE*
334 *Transactions on Information Theory*, 19(4):471–480, 1973.
- 335 Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. CommonsenseQA: A
336 question answering challenge targeting commonsense knowledge. In *Proceedings of the*
337 *2019 Conference of the North American Chapter of the Association for Computational Linguistics:*
338 *Human Language Technologies*, pp. 4149–4158. Association for Computational Linguistics,
339 2019. doi: 10.18653/v1/N19-1421. URL <https://aclanthology.org/N19-1421/>.
- 340 Johannes Welbl, Nelson F. Liu, and Matt Gardner. Crowdsourcing multiple choice science
341 questions, 2017. URL <https://arxiv.org/abs/1707.06209>.
- 342 Aaron D. Wyner and Jacob Ziv. The rate-distortion function for source coding with side
343 information at the decoder. *IEEE Transactions on Information Theory*, 22(1):1–10, 1976.

- 344 An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan
 345 Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report, 2025. URL
 346 <https://arxiv.org/abs/2412.15115>.
- 347 Amir Zandieh, Majid Daliri, and Insu Han. QJL: 1-bit quantized jl transform for kv cache
 348 quantization with zero overhead, 2024. URL <https://arxiv.org/abs/2406.03482>.
- 349 Amir Zandieh, Majid Daliri, Majid Hadian, and Vahab Mirrokni. TurboQuant: Online vector
 350 quantization with near-optimal distortion rate. In *International Conference on Learning*
 351 *Representations*, 2026. URL <https://openreview.net/forum?id=t03ASKZlok>.
- 352 Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. HellaSwag: Can
 353 a machine really finish your sentence? In *Proceedings of the 57th Annual Meeting of the*
 354 *Association for Computational Linguistics*, pp. 4791–4800. Association for Computational
 355 Linguistics, 2019. doi: 10.18653/v1/P19-1472. URL <https://aclanthology.org/P19-1472/>.
- 357 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao
 358 Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett,
 359 and Ying Sheng. SGLang: Efficient execution of structured language model pro-
 360 grams. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi:
 361 10.52202/079017-2000. URL [https://proceedings.neurips.cc/paper_files/paper/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html)
 362 [2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html).

363 A Reviewer-facing claim audit

364 Table 7 states the claim boundary used to write the main paper. The full generated packet is
 365 in results/latentwire_colm_v3_review_packet_20260505/.

Claim	Status	Evidence	Required wording
source-private packet protocol	supported	packet rows and controls	protocol/evaluation contri- bution
narrow same-family packet util- ity	supported but narrow	ARC-Challenge and Open- BookQA rows	source-candidate transfer
shortcut collapses are common	supported	source-index and wrong-row au- dits	claim boundary, not failure- only story
beats C2C or dense KV transfer	unsupported	no native matched C2C row	do not claim
native GPU/HBM/energy win	unsupported	no native NVIDIA measurement	future runbook only
broad latent/cross-family com- munication	unsupported	Phi-3 and connector failures	ICLR future goal

Table 7: Claim audit used for COLM_v3 integration. Each main-text claim must map to one supported or supported-but-narrow row.

366 B Artifact manifest

- 367 • ARC-Challenge headline: [colm_final/evidence/results/source_private_arc_](colm_final/evidence/results/source_private_arc_challenge_fourier_anchor_syndrome_gate_20260502_budget8_10seed.b2000)
 368 [challenge_fourier_anchor_syndrome_gate_20260502_budget8_10seed.b2000](colm_final/evidence/results/source_private_arc_challenge_fourier_anchor_syndrome_gate_20260502_budget8_10seed.b2000).
- 369 • Phi-3 falsification: [colm_final/evidence/results/source_private_arc_challenge_](colm_final/evidence/results/source_private_arc_challenge_source_family_cache_falsification_20260502_phi3_cpu_budget8_10seed.b2000)
 370 [source_family_cache_falsification_20260502_phi3_cpu_budget8_10seed.b2000](colm_final/evidence/results/source_private_arc_challenge_source_family_cache_falsification_20260502_phi3_cpu_budget8_10seed.b2000).
- 371 • Cross-family decomposition: [colm_final/evidence/results/source_private_arc_](colm_final/evidence/results/source_private_arc_cross_family_failure_decomposition_20260502)
 372 [cross_family_failure_decomposition_20260502](colm_final/evidence/results/source_private_arc_cross_family_failure_decomposition_20260502).
- 373 • Systems boundary: [colm_final/evidence/results/source_private_systems_](colm_final/evidence/results/source_private_systems_boundary_figure_table_20260502)
 374 [boundary_figure_table_20260502](colm_final/evidence/results/source_private_systems_boundary_figure_table_20260502).
- 375 • OpenBookQA headline: [colm_final/evidence/results/source_private_](colm_final/evidence/results/source_private_openbookqa_seed_stability_20260501_qwen05_hashed_test.3b)
 376 [openbookqa_seed_stability_20260501_qwen05_hashed_test.3b](colm_final/evidence/results/source_private_openbookqa_seed_stability_20260501_qwen05_hashed_test.3b).
- 377 • Source-index audit: [colm_final/evidence/results/source_private_colm_](colm_final/evidence/results/source_private_colm_acceptance_baselines_20260502)
 378 [acceptance_baselines_20260502](colm_final/evidence/results/source_private_colm_acceptance_baselines_20260502).
- 379 • COLM_v3 review packet: results/latentwire_colm_v3_review_packet_20260505.

C Reproducibility summary

The artifact bundle records exact shell commands, input hashes, output hashes, and rerun results. The paper rows are produced from cached, answer-key-forbidden source-choice artifacts with the following settings:

Row	script	seeds	key settings
ARC-Challenge	Fourier/anchor syndrome gate	10	8B, 2000 bootstraps
OpenBookQA	seed-stability gate	5	3B, hashed features, 500 bootstraps
Source-index audit	acceptance baseline audit	10/5	1000 bootstraps, 2/3/4/8B rate curve
Phi-3	source-family cache falsification	10	8B, cached Phi-3 source
Systems	systems boundary table	–	byte/exposure accounting only

Frozen artifacts are hash-verified. In the final local pass, OpenBookQA, systems, and failure-decomposition reruns completed; the more expensive full ARC-Challenge/Phi-3 reruns are represented by frozen hashes unless rerun offline. Rerun JSON files are metric-oriented rather than byte-identical because they contain creation times and some artifacts include measured local latencies.

Workshop checklist. The current paper is ready for reviewer circulation only after the generated COLM.v3 packet and the PDF agree on: (1) the main claim, (2) the source-private threat model, (3) all main table values, (4) measured-versus-estimated systems labels, and (5) unsupported claims. The repository command

```
./venv_arm64/bin/python scripts/build_latentwire_colm_v3_review_packet.py
```

regenerates the claim audit, table/figure inventory, systems ledger, artifact manifest, and NVIDIA native runbook from tracked source artifacts.

D Lay explanation of the main experiment

The experiment asks whether one model can send a tiny hint to another model. Both models see the same question and answer choices. The sender chooses a compressed coordinate hint, only a few bytes long. The receiver uses that hint to choose an answer. If both sides agree on the coordinate system, the hint helps. If we scramble the coordinate system, the hint stops helping. That is evidence for a real shared packet interface on the tested same-family setting, while the failed Phi-3 run shows the interface is not yet general across model families.