

Fixed-Byte Source-Private Candidate Packets for Bounded Model Collaboration

Anonymous authors

Paper under double-blind review

Abstract

Multi-model language systems usually communicate by emitting text or by sharing internal state such as KV caches. Text is portable but slow and lossy; cache sharing is information-rich but exposes thousands of bytes of model state per token and requires architecture-specific machinery. We study a narrower question: can a frozen source model send a fixed-byte, source-private candidate packet that improves a frozen receiver on public multiple-choice reasoning tasks, without sending source text, hidden states, KV cache, source logits, or gold labels? We present LatentWire, a public-coordinate packet protocol in which source and receiver agree on an anchor/spectral basis over candidate answers, the source emits a tiny signed packet, and the receiver decodes the packet against its own public side information. On ARC-Challenge test, an 8-byte packet reaches 0.344 accuracy across 10 seeds, above target-only 0.265 and same-budget structured text 0.300, with destructive coordinate controls failing in all 10 seeds; the source-choice accuracy is 0.346, so the positive should be read as compact source-candidate transfer rather than evidence synthesis beyond the source. On OpenBookQA test, a 3-byte packet reaches 0.378 across 5 seeds, above target-only 0.276 and same-budget text 0.350. We also report strict negative evidence: a Phi-3 cross-family source fails the same gate, and learned cache/score connectors on available TinyLlama/Qwen cached artifacts do not repair the broader cross-family gap. The contribution is therefore not a universal latent language claim and does not beat explicit source-index communication. It is a reproducible fixed-byte source-private transfer result, a falsification protocol for cross-family claims, and a systems accounting boundary showing that task packets are 69.8x smaller than even a one-token 1-bit-per-KV-element accounting floor.

1 Introduction

Future language-model systems will often use more than one model: a cheap model may screen candidates, a specialist may inspect hard cases, and a stronger receiver may combine evidence. The simplest interface is text. Recent cache-level systems instead communicate through model internals, such as fused or selectively shared KV caches (Fu et al., 2026; Shi et al., 2026). These approaches are powerful, but they couple communication to large hidden-state objects and serving substrates.

This paper asks whether a much smaller interface can carry useful source-side candidate evidence. In our setting, both models see the same public question and answer choices. The source performs private computation over those choices. The only communicated object is a short byte packet; the receiver never sees source text, source logits, source hidden vectors, KV cache, or the answer key. This is a side-information problem in the spirit of distributed source coding (Slepian & Wolf, 1973; Wyner & Ziv, 1976): the receiver already has the public prompt and its own computation, so the packet should carry only the source evidence that is useful at the decoder.

Our strongest current result is deliberately scoped. On ARC-Challenge and OpenBookQA (Clark et al., 2018; Mihaylov et al., 2018), same-family Qwen2.5 packet transfer (Yang et al., 2025) improves over target-only and same-budget text controls. However, a




Figure 1: LatentWire protocol. Both sides receive the public task. The source privately scores candidates and emits only a fixed-byte packet. The receiver decodes the packet with a public basis and its own side information. Destructive controls scramble the shared coordinate system or candidate relation.

strict Phi-3 source-family replacement (Abdin et al., 2024) fails, and additional connector repairs fail on the available TinyLlama/Qwen cached artifacts. This boundary matters: the paper is not claiming solved cross-family latent communication.

The contributions are:

- A fixed-byte, source-private packet protocol for multiple-choice candidate transfer with no source text, hidden vector, KV cache, source-logit, or gold-label exposure.
- A public-coordinate anchor/spectral implementation that is positive on ARC-Challenge and OpenBookQA under seed repeats, paired uncertainty, same-budget text controls, and destructive coordinate mismatch controls.
- A source-choice and falsification audit showing that the current positive is mostly source-choice preserving and same-family/source-cache specific: Phi-3 cross-family transfer and cached learned repair connectors fail.
- A systems accounting artifact comparing packet bytes and exposure classes against KV/cache communication and KV quantization floors, while explicitly separating byte accounting from native serving-speed claims.

2 Source-private packet transfer

Let a public multiple-choice instance be $x = (q, c_1, \dots, c_m)$ with answer choices c_i . A source model S and receiver R both observe x , but only S computes the source-side private evidence $e_S(x)$. A packet encoder emits

$$z = E(e_S(x), x), \quad |z| \leq B \text{ bytes.}$$

The receiver predicts

$$\hat{y} = D_R(x, z),$$

where D_R may use public task features and the receiver’s own state, but may not access source text, source hidden states, source KV cache, source logits, or answer labels.

We evaluate a packet only if it beats three surfaces: target-only $D_R(x, \emptyset)$, a same-budget structured text control, and destructive controls that preserve superficial packet size but destroy the intended source-receiver relation. We report seed repeats and paired bootstrap lower bounds, and we separately report source-choice accuracy where directly audited. This matters because a packet can be useful while mostly preserving the source’s selected candidate; such a result is source-private transfer, not proof of a shared latent language or of reasoning beyond the source endpoint.

Benchmark	split	bytes	seeds	source	packet	target	text	CI low
ARC-Challenge	test, 1172	8/11	10/10	0.346	0.344	0.265	0.300	+0.038
OpenBookQA	test, 500	3/6	5/5	–	0.378	0.276	0.350	+0.038

Table 1: Main positive rows. “Bytes” reports payload/framed bytes. “Source” is the directly audited, answer-key-forbidden source-choice accuracy before packet decoding; “–” means no separate source-choice row is recorded in the frozen artifact. “Text” is the same-budget structured text control. “CI low” is the minimum paired-bootstrap 95% lower bound for packet minus target-only.




Figure 2: Packet accuracy against target-only and same-budget structured text controls. Error bars show the lower paired-bootstrap lift above target-only; upper uncertainty is omitted for readability because the gate is one-sided.

74 3 Public-coordinate packets

75 The current ARC-Challenge implementation uses a public train-anchor chart. For each
76 candidate, we compute a public relative-coordinate feature against a frozen set of training
77 anchors, then project that feature into a low-frequency orthonormal DCT-II spectral basis.
78 The packet is a sparse signed syndrome in this shared coordinate system. The receiver
79 scores candidate compatibility with the packet in the same public basis, combining the
80 packet evidence with receiver-side candidate information.

81 The design is related to relative representations (Moschella et al., 2023): anchors provide a
82 coordinate system less tied to raw hidden dimensions. Our use is narrower. We do not claim
83 zero-shot latent stitching or global representation alignment; we use a public basis to define
84 a tiny task packet. Random shared anchors work nearly as well as semantic anchors in
85 the current hashed ARC-Challenge implementation, so the correct interpretation is shared
86 coordinate agreement, not semantic anchor names.

87 OpenBookQA uses the same packet-evaluation contract at a smaller 3-byte budget. Across
88 both tasks, the source-choice cache is answer-key-forbidden: the source sees public questions
89 and choices, not labels. The packet is evaluated on frozen test slices after validation selection.

90 4 Main results

91 Table 1 and Figure 2 summarize the current positive evidence. The ARC-Challenge row uses
92 10 packet/control seeds and 2000 paired bootstrap samples per seed on 1172 test examples.
93 The OpenBookQA row uses 5 seeds on the 500-example test slice.

Source	surface	n	packet	target	text	ref. pkt	blocker
Phi-3	full test	1172	0.244	0.265	0.232	–	source quality
Phi-3	Qwen-disagree	833	0.200	0.273	0.203	0.340	source quality
TinyLlama	Qwen-disagree	473	0.269	0.268	0.258	0.317	family mismatch
Qwen2.5-1.5B	full test	1172	0.442	0.265	0.401	–	val incomplete

Table 2: Cross-family and source-family diagnostics. “Ref. pkt” is the Qwen-substituted packet on disagreement rows. Qwen2.5-1.5B is encouraging but same-family and validation-incomplete, so it is not a headline cross-family claim.

On ARC-Challenge, the matched Fourier/anchor-syndrome packet reaches 0.344 mean test accuracy with minimum seed accuracy 0.342. The answer-key-forbidden source choice itself reaches 0.346, and packet predictions match the source-selected candidate at 0.995 mean rate across seeds. Anchor-ID shuffle, anchor-value shuffle, and spectral-bin permutation controls all fail in 0/10 seeds and collapse near target-only. The random shared-anchor diagnostic passes, which weakens any claim that named anchors are semantically special, but strengthens the claim that a shared public coordinate chart is the relevant mechanism.

On OpenBookQA, the 3-byte packet reaches 0.378–0.380 across seeds. The frozen artifact does not record a separate source-choice accuracy row, so this task is used as a second packet/control evidence row rather than as a source-choice audit. The largest candidate-derangement control accuracy is 0.212, and the minimum lift over same-budget structured text is +0.028. This second task matters because it reduces the risk that the result is only an ARC-Challenge artifact, though it is still a same-family packet result rather than a cross-family result.

5 Falsification and failure decomposition

The most important negative result is the Phi-3 source-family replacement. We keep the same ARC-Challenge packet budget, public basis, seed count, and bootstrap budget, but replace the Qwen source-choice cache with a Phi-3 source. The packet no longer passes. On the full ARC-Challenge test slice, the Phi-3 packet reaches 0.244 versus target-only 0.265. On the stricter Qwen-disagreement slice, where the alternate source and Qwen source disagree, the Phi-3 packet reaches 0.200 versus 0.340 for the Qwen-substituted packet.

The failure decomposition separates three causes: source endpoint quality, packet corruption, and common-feature mismatch. For Phi-3, packet predictions follow the Phi-3 source choice at 0.997, but the source itself is below the target. For TinyLlama, packet fidelity is also high, but the source-family mismatch causes the disagreement-slice packet to lose to Qwen-substituted packets. For Qwen2.5-1.5B, the source is stronger and the packet is useful on frozen test, but the validation disagreement gate is incomplete. Across these rows, the packet mostly transports the source-selected candidate through the public coordinate basis. The next research gate is therefore not another scalar packet shuffle; it is either an explicit comparison to source-index/score packets or a richer common-feature connector that can beat packet-only source choice.

We also tested cached learned repairs on the available local artifacts. A candidate-syndrome connector that sees packet and score-shape features reaches 0.288 versus 0.317 for Qwen-substituted packets on frozen TinyLlama/Qwen disagreement rows. A one-hidden-layer MLP over cached TinyLlama hidden/query summaries reaches 0.232, below Qwen-substituted packets, cached Tiny packets, target-only, and same-budget text. These negatives are useful: simple cached score geometry and mean hidden/query summaries do not explain the remaining oracle headroom.




Figure 3: Boundary-byte accounting. Packet rows are measured task-packet objects. KV/cache rows are lower-bound floors or neighboring serving substrates, not native speed comparisons.

Method	raw	framed	text	KV/hidden	claim status
LatentWire OpenBookQA packet	3B	6B	no	no	local positive
LatentWire ARC-Challenge packet	8B	11B	no	no	local positive
Same-byte text control	8B	11B	yes	no	control only
1-bit/KV-element floor	768B	768B	no	KV	floor only
KIVI 2-bit KV floor	1536B	1536B	no	KV	floor only
C2C one-token KV floor	12288B	12288B	no	KV	native pending

Table 3: Exposure classes. The paper-ready systems claim is that the communicated object is much smaller and less exposing than KV/cache-state objects under this accounting. It is not a GPU throughput, TTFT, TPOT, or HBM claim.

6 Systems boundary

The systems contribution is byte/exposure accounting, not native serving throughput. Figure 3 compares framed packet bytes with one-token KV/cache-state accounting floors from neighboring systems and quantization methods. LatentWire sends 6–11 framed bytes on the two headline rows. A one-token 1-bit-per-KV-element accounting floor is 768 bytes in our model-configuration accounting; the largest headline packet is therefore 69.8x smaller than that floor. C2C-style cache transfer, KVComm-style selective sharing, and serving substrates such as vLLM/PagedAttention and SGLang operate on KV/cache objects rather than task-level packets (Fu et al., 2026; Shi et al., 2026; Kwon et al., 2023; Zheng et al., 2024). KV quantization work reduces cache bytes but still communicates or stores model state (Zandieh et al., 2024; Liu et al., 2024; Hooper et al., 2024; Zandieh et al., 2026).

7 Related work

Latent and relative representation communication. Relative representations use similarities to anchors to enable latent-space communication under invariances (Moschella et al., 2023). Sparse autoencoder work motivates interpretable feature coordinates (Huben et al., 2024), while feature-space universality work suggests that some features can be related across models (Lan et al., 2024). SAE Bench provides evaluation machinery for sparse au-

149 toencoders (Karvonen et al., 2025), but these works do not by themselves provide a packet
150 protocol or task-level evidence transfer result.

151 **Prompt compression and learned prefixes.** Prefix tuning and gist tokens learn continuous
152 or discrete prompt-like interfaces within language models (Li & Liang, 2021; Mu et al., 2023).
153 LatentWire differs because the communicated object is a byte packet decoded by a public
154 receiver, not a learned soft prompt injected into one target. Query-bottleneck connectors
155 such as BLIP-2’s Q-Former show that learned bottlenecks can bridge frozen model families
156 across modalities (Li et al., 2023); our failed cached hidden/query summaries suggest that a
157 true tokenwise query-bottleneck is the right future branch, but it is not closed here.

158 **Cache communication and systems.** Cache-to-cache and selective KV-sharing systems
159 communicate richer internal state and are the closest competitors for model-to-model
160 communication (Fu et al., 2026; Shi et al., 2026). vLLM and SGLang improve serving
161 through KV-cache management and reuse (Kwon et al., 2023; Zheng et al., 2024). KV
162 quantizers such as QJL, KIVI, KVQuant, and TurboQuant reduce state cost but remain
163 state-transfer or state-storage techniques (Zandieh et al., 2024; Liu et al., 2024; Hooper et al.,
164 2024; Zandieh et al., 2026). Our current systems result is a boundary artifact against these
165 state objects, not a native throughput win.

166 8 Limitations and next gate

167 The main limitation is scope. The positive results are same-family/source-cache results
168 on public multiple-choice tasks, and the packet largely preserves the source’s selected
169 candidate. We therefore do not claim to beat explicit source-index communication, source-
170 score quantization, or calibrated source-choice baselines; those are required before a stronger
171 compression claim. Phi-3 cross-family replacement fails, TinyLlama common-feature repair
172 fails, and a true cross-family learned query-bottleneck remains future work. We also do not
173 claim superiority over C2C, KVComm, QJL, KIVI, TurboQuant, vLLM, or SGLang in native
174 serving metrics. Those rows require NVIDIA hardware and matched implementations.

175 The next exact gate for an ICLR-strength claim is a frozen-endpoint tokenwise connector with
176 an explicit byte or query budget, tested against C2C/KV-style and structured-text baselines,
177 with source-destroying controls and strict same-family versus cross-family separation.
178 The current strongest claim is narrower: fixed-byte source-private packets can transfer
179 useful same-family evidence on ARC-Challenge and OpenBookQA, and the accompanying
180 falsification ladder clarifies where the method stops working.

181 9 Conclusion

182 LatentWire shows that a tiny source-private candidate packet can improve frozen receiver
183 decisions on two public QA benchmarks under strict same-family controls. The same evi-
184 dence also draws a hard boundary: the current method is mostly source-choice preserving,
185 is not a universal latent language, and does not yet solve cross-family communication. This
186 combination of positive packet transfer, negative falsification, and systems byte/exposure
187 accounting is the paper’s core.

188 References

- 189 Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan,
190 Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3
191 technical report: A highly capable language model locally on your phone, 2024. URL
192 <https://arxiv.org/abs/2404.14219>.
- 193 Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick,
194 and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning
195 challenge, 2018. URL <https://arxiv.org/abs/1803.05457>.

- 196 Tianyu Fu, Zihan Min, Hanling Zhang, Jichao Yan, Guohao Dai, Wanli Ouyang, and
197 Yu Wang. Cache-to-cache: Direct semantic communication between large language
198 models. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=LeatkxrBCi>.
199
- 200 Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia
201 Shao, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 million con-
202 text length llm inference with kv cache quantization. In *Advances in Neural*
203 *Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-0040.
204 URL [https://proceedings.neurips.cc/paper_files/paper/2024/hash/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/028fcbcf85435d39a40c4d61b42c99a4-Abstract-Conference.html)
205 [028fcbcf85435d39a40c4d61b42c99a4-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/028fcbcf85435d39a40c4d61b42c99a4-Abstract-Conference.html).
- 206 Robert Huben, Hoagy Cunningham, Logan Riggs Smith, Aidan Ewart, and Lee Sharkey.
207 Sparse autoencoders find highly interpretable features in language models. In *International*
208 *Conference on Learning Representations*, 2024. URL [https://openreview.net/forum?id=](https://openreview.net/forum?id=F76bwRSLeK)
209 [F76bwRSLeK](https://openreview.net/forum?id=F76bwRSLeK).
- 210 Adam Karvonen, Can Rager, Johnny Lin, Curt Tigges, Joseph Isaac Bloom, David Chantin,
211 Yeu-Tong Lau, Eoin Farrell, Callum Stuart Mcdougall, Kola Ayonrinde, Demian Till,
212 Matthew Wearden, Arthur Conmy, Samuel Marks, and Neel Nanda. SAEBench: A
213 comprehensive benchmark for sparse autoencoders in language model interpretability.
214 In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of
215 *Proceedings of Machine Learning Research*, pp. 29223–29264. PMLR, 2025. URL <https://proceedings.mlr.press/v267/karvonen25a.html>.
216
- 217 Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu,
218 Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large
219 language model serving with pagedattention. In *Proceedings of the 29th Symposium on*
220 *Operating Systems Principles*, pp. 611–626, 2023. doi: 10.1145/3600006.3613165. URL
221 <https://dl.acm.org/doi/10.1145/3600006.3613165>.
- 222 Michael Lan, Philip Torr, Austin Meek, Ashkan Khakzar, David Krueger, and Fazl Barez.
223 Quantifying feature space universality across large language models via sparse autoen-
224 coders, 2024. URL <https://arxiv.org/abs/2410.06981>.
- 225 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. BLIP-2: Bootstrapping language-
226 image pre-training with frozen image encoders and large language models. In *Proceedings*
227 *of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine*
228 *Learning Research*, pp. 19730–19742. PMLR, 2023. URL [https://proceedings.mlr.press/](https://proceedings.mlr.press/v202/li23q.html)
229 [v202/li23q.html](https://proceedings.mlr.press/v202/li23q.html).
- 230 Xiang Lisa Li and Percy Liang. Prefix-tuning: Optimizing continuous prompts for generation.
231 In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and*
232 *the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*,
233 pp. 4582–4597. Association for Computational Linguistics, 2021. doi: 10.18653/v1/2021.
234 [acl-long.353](https://aclanthology.org/2021.acl-long.353/). URL <https://aclanthology.org/2021.acl-long.353/>.
- 235 Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhuo Xu, Vladimir Braverman,
236 Beidi Chen, and Xia Hu. KIVI: A tuning-free asymmetric 2bit quantization for kv cache.
237 In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of
238 *Proceedings of Machine Learning Research*, pp. 32332–32344. PMLR, 2024. URL <https://proceedings.mlr.press/v235/liu24bz.html>.
239
- 240 Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor
241 conduct electricity? a new dataset for open book question answering. In *Proceedings of*
242 *the 2018 Conference on Empirical Methods in Natural Language Processing*, pp. 2381–2391.
243 Association for Computational Linguistics, 2018. doi: 10.18653/v1/D18-1260. URL
244 <https://aclanthology.org/D18-1260/>.
- 245 Luca Moschella, Valentino Maiorca, Marco Fumero, Antonio Norelli, Francesco Locatello,
246 and Emanuele Rodolà. Relative representations enable zero-shot latent space com-
247 munication. In *International Conference on Learning Representations*, 2023. URL [https://openreview.net/forum?id=](https://openreview.net/forum?id=Src-nwieGJ)
248 [Src-nwieGJ](https://openreview.net/forum?id=Src-nwieGJ).

- 249 Jesse Mu, Xiang Li, and Noah Goodman. Learning to compress prompts with
 250 gist tokens. In *Advances in Neural Information Processing Systems*, volume 36,
 251 2023. URL [https://proceedings.neurips.cc/paper_files/paper/2023/hash/](https://proceedings.neurips.cc/paper_files/paper/2023/hash/3d77c6dcc7f143aa2154e7f4d5e22d68-Abstract-Conference.html)
 252 [3d77c6dcc7f143aa2154e7f4d5e22d68-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/3d77c6dcc7f143aa2154e7f4d5e22d68-Abstract-Conference.html).
- 253 Xiangyu Shi, Marco Chiesa, Gerald Q. Maguire Jr., and Dejan Kostic. KVComm: Enabling
 254 efficient llm communication through selective kv sharing. In *International Conference on*
 255 *Learning Representations*, 2026. URL <https://openreview.net/forum?id=F7rUng23nw>.
- 256 David Slepian and Jack K. Wolf. Noiseless coding of correlated information sources. *IEEE*
 257 *Transactions on Information Theory*, 19(4):471–480, 1973.
- 258 Aaron D. Wyner and Jacob Ziv. The rate-distortion function for source coding with side
 259 information at the decoder. *IEEE Transactions on Information Theory*, 22(1):1–10, 1976.
- 260 An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan
 261 Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report, 2025. URL
 262 <https://arxiv.org/abs/2412.15115>.
- 263 Amir Zandieh, Majid Daliri, and Insu Han. QJL: 1-bit quantized jl transform for kv cache
 264 quantization with zero overhead, 2024. URL <https://arxiv.org/abs/2406.03482>.
- 265 Amir Zandieh, Majid Daliri, Majid Hadian, and Vahab Mirrokni. TurboQuant: Online vector
 266 quantization with near-optimal distortion rate. In *International Conference on Learning*
 267 *Representations*, 2026. URL <https://openreview.net/forum?id=t03ASKZlok>.
- 268 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Sun, Jeff Huang, Cody Hao
 269 Yu, Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett,
 270 and Ying Sheng. SGLang: Efficient execution of structured language model pro-
 271 grams. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi:
 272 10.52202/079017-2000. URL [https://proceedings.neurips.cc/paper_files/paper/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html)
 273 [2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/724be4472168f31ba1c9ac630f15dec8-Abstract-Conference.html).

274 A Artifact manifest

- 275 • ARC-Challenge headline: [colm_final/evidence/results/source_private_arc_](#)
 276 [challenge-fourier_anchor_syndrome_gate_20260502_budget8_10seed.b2000](#).
- 277 • Phi-3 falsification: [colm_final/evidence/results/source_private_arc_challenge_](#)
 278 [source_family_cache_falsification_20260502_phi3_cpu_budget8_10seed.b2000](#).
- 279 • Cross-family decomposition: [colm_final/evidence/results/source_private_arc_](#)
 280 [cross_family_failure_decomposition_20260502](#).
- 281 • Systems boundary: [colm_final/evidence/results/source_private_systems_](#)
 282 [boundary_figure_table_20260502](#).
- 283 • OpenBookQA headline: [colm_final/evidence/results/source_private_](#)
 284 [openbookqa_seed_stability_20260501_qwen05_hashed.test.3b](#).

285 B Reproducibility summary

286 The artifact bundle records exact shell commands, input hashes, output hashes, and rerun
 287 results. The paper rows are produced from cached, answer-key-forbidden source-choice
 288 artifacts with the following settings:

Row	script	seeds	key settings
289 ARC-Challenge	Fourier/anchor syndrome gate	10	8B, 2000 bootstraps
OpenBookQA	seed-stability gate	5	3B, hashed features, 500 bootstraps
Phi-3	source-family cache falsification	10	8B, cached Phi-3 source
Systems	systems boundary table	–	byte/exposure accounting only

290 Frozen artifacts are hash-verified. In the final local pass, OpenBookQA, systems, and failure-
291 decomposition reruns completed; the more expensive full ARC-Challenge/Phi-3 reruns
292 are represented by frozen hashes unless rerun offline. Rerun JSON files are metric-oriented
293 rather than byte-identical because they contain creation times and some artifacts include
294 measured local latencies.

295 **C Lay explanation of the main experiment**

296 The experiment asks whether one model can send a tiny hint to another model. Both models
297 see the same question and answer choices. The sender chooses a compressed coordinate hint,
298 only a few bytes long. The receiver uses that hint to choose an answer. If both sides agree on
299 the coordinate system, the hint helps. If we scramble the coordinate system, the hint stops
300 helping. That is evidence for a real shared packet interface on the tested same-family setting,
301 while the failed Phi-3 run shows the interface is not yet general across model families.