

# HybridKernel: A Profiler-Driven Boundary-Fusion Gate for Hybrid Attention/SSM Serving

Anonymous authors

Paper under double-blind review

## Abstract

Hybrid attention/state-space language models invite a systems question: do attention-to-SSM layer boundaries create avoidable conversion, materialization, launch, or locality overhead that a fused boundary operator could remove? We report a Mac-feasible pre-GPU gate for this question. Architecture maps show large boundary-crossing activation streams, but a runtime audit weakens the broad novelty claim because current vLLM hybrid SSM support already handles important state-layout and transfer paths. A threshold model says Granite 4.0 H needs roughly 25% genuinely avoidable boundary traffic at 60% recovery to clear a 3% proxy gain; Qwen3-Next needs 10.4% but is less directly matched to the Granite Mamba2 hypothesis. We therefore do not claim a kernel speedup. The contribution is a reproducible profiler gate, a fixed-request vLLM driver, a native packet skeleton generator, a pre-registered profiler-analysis parser, and an artifact-completeness verifier that rejects client-only traces, unfilled skeletons, and stale analysis before any native result can be cited.

## 1 Hypothesis

HybridKernel tests whether adjacent attention and SSM layers expose boundary-local overhead not already removed by serving systems. Candidate costs include layout conversion, residual-stream materialization, host launch gaps, cache-locality loss, or redundant memory traffic. The hypothesis is deliberately narrow: a useful kernel must compose with existing vLLM hybrid mechanisms rather than replace them.

## 2 Current Evidence

## 3 Profiler Gate

The next experiment is a native NVIDIA/vLLM run with Nsight Systems and Nsight Compute. It must show a distinct boundary-local overhead attributable to attention/SSM transitions, not warmup, graph capture, batching, or unrelated kernels. The runbook now specifies exactly how to reduce Nsight traces into parser fields: total step time, boundary-local time, matched non-boundary control time, and recoverable fraction. Promotion requires at least three repeated runs whose recoverable-gain upper bound clears 3%. The analysis parser computes

$$\text{gain}_{ub} = \frac{\max(\text{boundary ms} - \text{matched control ms}, 0)}{\text{total step ms}} \times \text{recoverable fraction}.$$

If repeated native summaries show less than 1% recoverable gain, the branch should be killed or shelved. Before any run is cited, the artifact verifier must also pass on the exact run directory, confirming environment metadata, server-side Nsight Systems and Nsight Compute scope, profiler artifacts, profiling logs, the pre-registered readout questions, at least three distinct repeated metric rows for one model, and profiler-analysis outputs that match the same metric file. A packet generator creates this directory shape for GPU operators, but the checker rejects its `TODO_NATIVE_PROFILE_FILL` sentinels until real metadata,

| Artifact                 | Result                                                                                                                                                                                                      | Decision                                                                                 |
|--------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------|
| Architecture map         | Granite has 8 boundaries and 20.0% boundary stream fraction; Qwen3-Next has 23 inferred boundaries and 47.9%.                                                                                               | Inspectable, not speed evidence.                                                         |
| Runtime audit            | vLLM already handles hybrid state layout and transfer paths.                                                                                                                                                | Broad novelty weakened.                                                                  |
| Threshold model          | Granite needs about 25.0% avoidable boundary traffic at 60% recovery for a 3% proxy gain.                                                                                                                   | Mac kernel work not justified.                                                           |
| Profiler driver          | Fixed-request OpenAI-compatible driver dry-runs locally; runbook now profiles the vLLM server, not just the HTTP client.                                                                                    | Native gate is less likely to be mis-run.                                                |
| Analysis parser          | JSON profiler summaries are reduced into avoidable-boundary share and recoverable-gain upper bound.                                                                                                         | Decision rule is pre-registered.                                                         |
| Artifact verifier        | Native run directories are checked for metadata, server-side Nsight Systems and Compute scope, Nsight artifacts, logs, readout rows, distinct repeated metric rows, and matching profiler-analysis outputs. | Client-only, duplicated-trace, and stale-analysis evidence is rejected.                  |
| Packet generator         | A local script creates the required native run-packet skeleton with TODO sentinels.                                                                                                                         | GPU operators get a consistent handoff; skeletons are rejected as evidence until filled. |
| Synthetic packet fixture | A Mac-local fixture documents the packet shape and exercises the checker with placeholder Nsight files.                                                                                                     | Fixture is not profiler evidence.                                                        |
| Packet checklist         | The native handoff now has a single required packet layout and stop-local-work decision.                                                                                                                    | Local readiness is saturated until GPU data exists.                                      |

Table 1: HybridKernel evidence before native GPU profiling.

metrics, and readout entries replace them. This is an admissibility check only; it does not promote the method.

## 4 Claim Boundary

Allowed now: HybridKernel is a weakly alive profiler-driven systems branch with a runnable native measurement plan. Not allowed: throughput, latency, HBM, energy, occupancy, or vLLM speedup claims. If native profiling does not reveal separable boundary overhead, the branch should be killed.

## 5 Artifacts

- experimental/hybridkernel/phase2/nvidia\_vllm\_profiler\_runbook.md
- experimental/hybridkernel/phase2/profiler\_driver.py
- experimental/hybridkernel/phase2/create\_native\_run\_packet.py
- experimental/hybridkernel/phase2/analyze\_profiler\_metrics.py
- experimental/hybridkernel/phase2/check\_profiler\_run\_artifacts.py
- experimental/hybridkernel/phase2/native\_run\_packet\_checklist.md
- experimental/hybridkernel/phase2/profiler\_analysis\_gate.md

- `experimental/hybridkernel/phase2/pre_gpu_threshold_model.md`
- `experimental/hybridkernel/paper/colm_workshop_scaffold.md`