

SSQ-LR: Sub-FP16 Recurrent State Quantization for Hybrid Reasoners

Anonymous authors

Paper under double-blind review

Abstract

This is a living paper shell for a preregistered Mac-first experiment. SSQ-LR tests whether recurrent SSM state in hybrid reasoning models can be quantized below FP16 while preserving quality. No positive result or native systems claim is made yet.

1 Current Gate

The branch must pass state-distribution heterogeneity, state-only quantization sensitivity, and cross-model transfer gates before any GPU validation.