

ThoughtFlow-FP8: A Mixed Gate for Phase-Aware KV Retention

Anonymous authors

Paper under double-blind review

Abstract

ThoughtFlow-FP8 tests whether explicit anchor/phase retention plus byte-aware FP8-style budgeting can preserve useful reasoning context better than matched-budget KV-retention proxies. Synthetic traces support the intuition that phase markers and anchors can be protected, but real retained-context quality weakens the branch. On 24 distilgpt2 traces at a 0.20 keep fraction, a bounded successor policy, ThoughtFlow-saliency-recent, improves over the original phase-only policy and beats LongFlow-like and ThinKV-like proxies, but narrowly loses to an R-KV-like retained-prefix proxy on continuation NLL (3.434 vs 3.419). A held-out sweep over recency and saliency weights then finds a matched-budget ThoughtFlow-family policy that ties R-KV-like on held-out traces (3.480 vs 3.482 NLL; positive margin +0.001). We therefore present this as a mixed workshop gate, not a positive method. The next useful step is real hidden/KV saliency telemetry; GPU sparse-KV work is not justified until that gate clears.

1 Policy Family

The original policy protects anchors and phase labels. The first successor adds a small recent-token reserve. The second successor protects anchors, reserves half the retained budget for recent tokens, and lets phase markers compete with math-state and high-importance reasoning tokens under a saliency bonus. This tests whether the failure was caused by over-preserving interpretable markers that do not matter for immediate continuation likelihood.

Policy	Keep rate	NLL	Δ NLL vs full	PPL
full context	1.000	2.101	0.000	9.7
R-KV-like	0.210	3.419	1.319	38.7
ThinKV-like	0.210	3.583	1.482	44.6
LongFlow-like	0.210	3.961	1.861	74.2
ThoughtFlow	0.210	3.961	1.861	74.2
ThoughtFlow-recent	0.210	3.562	1.461	44.2
ThoughtFlow-saliency-recent	0.210	3.434	1.333	38.1
Held-out R-KV-like	0.211	3.482	–	43.7
Held-out sweep best	0.211	3.480	–	43.7

Table 1: Retained-context continuation proxy on distilgpt2 traces. Lower NLL is better. Full context is not a matched-budget baseline. The held-out sweep row is selected on 12 train traces and reported on 12 held-out traces, so its NLL is not directly averaged with the all-trace rows.

2 Decision

The branch is mixed but not revived as a positive method. ThoughtFlow-saliency-recent repairs most of the phase-only failure, and a small held-out sweep reaches a tie-range result against R-KV-like at matched budget. This is still text-prefix scoring, not real cache

dropping. A paper-ready claim would require a hidden/KV saliency policy that beats R-KV-like, LongFlow-like, and ThinKV-like rows at matched bytes, then validates under real sparse-KV or cache-dropping execution.

3 Artifacts

- `experimental/thoughtflow_fp8/phase2/perplexity_impact_proxy.md`
- `experimental/thoughtflow_fp8/phase2/policy_sweep.md`
- `experimental/thoughtflow_fp8/phase2/hidden_saliency_retention_probe.md`
- `experimental/thoughtflow_fp8/phase2/real_trace_retention_sweep.md`
- `experimental/thoughtflow_fp8/paper/colm_workshop_shell.md`