

When Byte-Scale Model Communication Should Be Text

Abstract

LatentWire asks whether a source model can send a small no-text packet that improves a receiver model beyond ordinary score and text baselines. The completed evidence supports a bounded negative for the discrete, byte-scale regime that is currently claim-ready. One-way score and residual packets do not beat source-index/confidence controls on the held-out aggregate, and the powered score ladder shows why: apparent source information drops from 2.098527 bits to 0.281933 bits after conditioning on receiver evidence. The strongest discrete-evidence escape fails in the opposite direction: when the same exact symbolic evidence is exposed as a visible 2-byte public code, it reaches 1.000 accuracy versus 0.775 for the matched learned packet, a +0.225 advantage with 95% CI [0.192188, 0.257812]. The last cheap privacy-bottleneck escape also fails: L-IB1 is KILL_UTILITY_IS_IDENTITY, with the current packet at utility/leakage 0.875/1.000 and the best bottleneck at only 0.602339/0.602339. The contribution is a claim-boundary result: for discrete low-byte evidence, no-text latent packets are dominated by score controls or equal-byte visible evidence unless the method moves to continuous dense state or a genuinely privacy-preserving bottleneck.

Claim-Boundary Box

Supported	Unsupported	Future-only
Score-like source packets are bounded negative under source-index/confidence controls.	A deployable-positive claim is unsupported.	Continuous dense cache/state transfer may still work, but it is outside this byte-scale discrete packet result.
Exact discrete evidence is text-capturable under the same byte budget in the local cache.	L-IB1 is not a privacy positive; it remains KILL_UTILITY_IS_IDENTITY.	A reviewed neural privacy bottleneck would need matched visible/anonymized text controls.

Supported	Unsupported	Future-only
Gold-aware oracle ceilings are rejected as leakage or gap-to-perfect measurements.	No cross-family, dense-cache, native GPU, or systems-speed claim is made.	Stronger generated-candidate reranking remains possible only with gold-blind verifier/source/target scores.

1. Introduction

Small model-to-model messages are attractive because they promise communication without exposing full reasoning traces. The hard question is not whether a helper model knows something useful. It is whether a byte-scale no-text packet conveys information that the receiver cannot already recover from its own scores, the source identity, or an equal-byte visible code.

This campaign originally searched for a positive LatentWire method. The result is a stronger and more useful boundary: the discrete evidence that makes a packet useful is visible-code capturable, while score-like residual signals mostly disappear after receiver conditioning. The paper therefore argues for a conservative standard for future model communication work: the relevant delta is not improvement over target-only, but improvement beyond source-index, source-confidence, equal-byte text, wrong-row, and destructive controls.

Contributions:

1. A controlled bounded negative for deployable score/residual packets. On the held-out aggregate with 381 rows, the deployable score packet is -0.023622 below source-index/confidence with CI $[-0.060367, 0.013123]$. On the powered dev/gate ladder with 1500 scored rows and 359 gate rows, the current deployable packet is -0.013928 below the best score baseline with CI $[-0.050139, 0.022284]$.
2. A mechanism for the negative result. Source scores contain 2.098527 bits before receiver conditioning, but only 0.281933 bits after conditioning on receiver evidence (Figure 1).
3. A direct exact-discrete evidence test. On 640 model-helper rows over 160 unique examples, a matched model packet reaches 0.775 accuracy, while the exact same symbolic evidence serialized as visible public code reaches 1.000 (Figure 2). Random, shuffled, answer-only, and target-only controls remain at 0.250.
4. A final cheap privacy escape test. L-IB1 cannot preserve current-packet utility while hiding source/evidence identity (Figure 3).




Figure 1: Receiver conditioning bits

Figure 1: The source-score signal largely disappears once receiver evidence is included.

2. Experimental Design

All results in this paper are drawn from existing aggregate or cached non-claim artifacts. No new experiment, GPU run, MPS live forward, confirmation split access, or killed-method rerun is used in this paper-finalization pass. The package includes generated figures, provenance tables, and review records only.

The main evidence families are:

- deployable one-way score/residual packets;
- the powered score ladder and receiver-conditioning information estimate;
- L-Q1 receiver-query packet controls;
- exact-discrete evidence versus equal-byte visible code;
- L-IB1 cached privacy-bottleneck frontier;
- rejected oracle ceilings and dense/cache smokes used only as boundary evidence.

The paper’s numeric claims map to package-local figures and tables. The detailed provenance is in `tables/provenance.md`; the falsification ladder is in `tables/falsification_ladder.md`.

3. Falsification Ladder

The central asset of this paper is not a single negative row. It is a ladder of increasingly plausible escapes, each killed or parked by a concrete hard baseline.

See `tables/falsification_ladder.md` for the complete table. The short version is:

Escape	Result	Controlling baseline
Score/WZ one-way packet	held-out delta -0.023622, CI [-0.060367, 0.013123]	source- index+confidence
Powered score ladder	current packet delta -0.013928, CI [-0.050139, 0.022284]	best equal-byte score/source baseline
L-B1 trust packet	killed as source-copy/damage- avoidance failure	source-index/source- selected metadata
L-Q1 receiver query	delta -0.022284, CI [-0.055710, 0.011142]	source- index+confidence and ablations
Exact discrete evidence	visible code beats packet by +0.225	equal-byte public code
L-IB1 privacy bottleneck	best bottleneck 0.602339/0.602339 vs current 0.875/1.000	same-byte code and anonymized text
Oracle fuser/verifier ceilings	rejected as gold-aware or setup-blocked	gold-leakage audit

4. Results

4.1 Deployable Packets Do Not Beat Source-Index Confidence

Figure 2: On the held-out aggregate, the deployable packet does not beat the source-index/confidence baseline. The source+target-at-encoder upper bound is positive but non-deployable.

The held-out aggregate shows the exact boundary. Target-only accuracy is 0.194226, source-index+confidence is 0.204724, and the deployable WZ packet is 0.181102. The source+target-at-encoder upper bound reaches 0.314961, so the experiment is not merely noise. The problem is deployability: the useful




Figure 2: Held-out source-index null

signal appears when the encoder has target evidence, not in the allowed source-only low-byte packet.

4.2 Receiver Conditioning Explains The Null

The powered ladder estimates 2.098527 bits of source-score information before receiver conditioning and 0.281933 bits afterward. Same-family models on shared multiple-choice tasks often fail and succeed for overlapping reasons. A source packet can therefore appear useful until the receiver’s own evidence and the source identity are included as baselines.

This result also explains why source-only oracle gaps are not enough. A positive source+target-at-encoder upper bound says there is information in the joint state; it does not say a deployable source-only byte packet transmits it.

4.3 Exact Discrete Evidence Is Better Sent As Text

Figure 3: The exact visible public signature reaches 1.000 accuracy, while the learned matched packet reaches 0.775. Controls remain at chance.

The exact-discrete evidence test is the cleanest negative result. The matched learned packet is useful, but the exact symbolic evidence that makes it useful is also public-code serializable inside the same byte budget. When the receiver gets that code visibly, it reaches perfect accuracy in this cache.

This is not a universal impossibility theorem. The sample has 640 model-helper rows but only 160 unique examples. We therefore state it as an operational




Figure 3: Exact discrete evidence

theorem for this class of evidence: if the useful content is a finite symbolic evidence signature and equal-byte visible public code is allowed, the opaque packet has no inherent no-text advantage.

4.4 L-IB1 Does Not Produce A Privacy-Positive Packet

Figure 4: L-IB1 cannot move off the utility/leakage identity diagonal enough to preserve utility while hiding source/evidence structure.

L-IB1 asked whether an adversarial feature selector could keep utility while hiding source/evidence identity on 512 cached rows. It could not. The current high-utility packet has cached utility 0.875 and max leakage 1.000; the best bottleneck has proxy utility 0.602339 and max leakage 0.602339. The exact public code matches current-packet utility and leakage, while adaptive anonymized text remains a necessary comparator for any future privacy claim.

The verdict remains KILL_UTILITY_IS_IDENTITY. This is a cached gate, not a formal privacy impossibility result. It kills the cheap local escape and sets the bar for a future build-scale IB method.

4.5 Rejected Oracle Ceilings Are Not Communication Claims

Two large apparent ceilings are explicitly excluded. L-PC5's strict rerank oracle and L-C2's strict fuser oracle measure answer-aware or setup-blocked gaps, not deployable no-text communication. The gold-blind L-PC5 attempt reaches only




Figure 4: L-IB1 utility leakage frontier

82 prompts, below its required floor, and ties the equal-byte text control. L-C2 has no reviewed gold-free hidden-state fuser objective locally.

These results are useful only as warnings. A high oracle ceiling does not imply a latent packet win unless the scorer/fuser is gold-blind and beats equal-byte text/source-index controls.

5. Related Work

Continuous cache/state communication systems such as C2C, KVComm, Interlat, and Latent Cache Flow occupy a different lane: they transmit or align dense hidden state rather than a small discrete evidence packet. That lane remains a plausible future route but is not claimed here.

Privacy-preserving semantic communication and adaptive text anonymization motivate the privacy frontier. In this paper they function as baselines and scope limits: a no-text privacy claim must beat adaptive visible text at the same byte budget while demonstrating lower leakage.

The closest practical baselines for this work are deliberately simple: source identity, source confidence, source rank, equal-byte score sketches, equal-byte visible evidence, wrong-row packets, and shuffled packets. These baselines are the point. They turn raw helper gains into deployable communication claims only when the gain survives.

6. Limitations

The exact-discrete evidence test has 640 model-helper rows but only 160 unique examples. The result is strong for this cache and evidence class, but it is not a theorem over all protocols.

L-IB1 is a CPU cached feature-selector test. It does not train a live neural encoder, run new generation, or test continuous dense state.

The held-out one-way result is used as aggregate evidence only. Raw held-out rows are not packaged, and no paper-finalization step reads or reruns them.

The paper does not claim that latent communication is impossible. It claims that the completed byte-scale discrete packet regimes do not establish a no-text advantage over hard visible/score baselines.

7. Conclusion

LatentWire did not produce a deployable no-text byte-scale positive. It produced a sharp boundary. Score-like source packets are mostly redundant after receiver conditioning, exact discrete evidence is better transmitted as equal-byte

visible public code, and the cheap privacy bottleneck cannot separate utility from identity leakage. Future positives must leave this killed regime: continuous dense state, or a real privacy-preserving bottleneck that beats adaptive same-byte text.