
Channel-Set Drift in Long-Reasoning W4A16 LLMs: Mechanisms and a Budget-Tuned Remedy

Anonymous authors

Paper under double-blind review

Abstract

Static per-channel protection is attractive for W4A16 reasoning inference, but it assumes that high-magnitude activation channels remain useful protection targets throughout decode. We measure decode-time channel-set drift across four reasoning LLMs and find strict top-1% set-leaving of 53–67% by horizons up to 20K tokens. The effect appears at the measured block-output surface in hybrid Mamba-2, parallel-hybrid, and pure-Transformer models, extending DecDEC’s short-horizon Transformer observation and contradicting the channel-persistence assumption behind Quamba2-style static ordering in our long-decode setting. We then test seven channel-set protection families with matched-cost and random controls. Hard position or scale switches are actively harmful, top-1% smoothing and cross-tensor coupling have signal but insufficient recovery, and reactive DecDEC-style selection does not recover the W4A16 gap. Budget-tuned EMA protection is the strongest channel-set remedy: on Nemotron-3, top-10 protection recovers 0.82 of the BF16-vs-static gap and beats static top-10 by 0.22; on Granite-Small, top-5 recovers 0.45 with a wide CI. A ParoQuant rotation baseline recovers 0.75 on Granite-Small, scoping our negative results to channel-set protection rather than W4A16 PTQ generally. KL trajectories grow sublinearly with AR decays 0.44–0.51, weakening compound-error explanations in the measured regime.

1 Introduction

Long-reasoning models can spend thousands of decode steps in a single answer, so calibration-time quantization decisions must survive within-trace distribution shift. Static protected-channel maps are especially brittle if the channels that dominate early activation magnitude do not remain dominant later. This paper studies that assumption for W4A16 post-training quantization (PTQ): 4-bit weights, 16-bit activations, and a small protected activation-channel set kept at higher precision.

The assumption matters for recent state-space quantization. Quamba2 argues that Mamba-family activations exhibit channel-order preservation and activation persistence, properties that support static channel ordering and clustering (Chiang et al., 2025). We test the corresponding protected-set assumption at long reasoning horizons. Across Granite-4.0-H-Small, Nemotron-3-Nano, DeepSeek-R1-Distill-Qwen-1.5B, and Falcon-H1, 53–67% of the top-1% channels at the calibration position strictly leave the later top-1% set. This is not merely within-set rank shuffling; it is membership change in the set a static policy would protect.

The systems motivation is narrow but practical. Reasoning workloads are expensive, and inference-cost forecasts expect large pressure to reduce per-token cost even as reasoning demand grows (Bort, 2025; Gartner, 2026). Recent W4A16 systems for larger standard Transformers can recover substantial quality, including ParoQuant’s rotation-based reasoning results (Liang et al., 2025). Our focus is different: small reasoning models, hybrid Mamba-2 or parallel-hybrid architectures, pure PTQ rather than distillation, and channel-set protection under long decode.

Our contributions are fourfold.

- Empirical drift. We provide a four-model cross-architecture measurement of decode-time channel-set drift up to 20K tokens, with strict set-leaving of 53–67% across hybrid Mamba-2, parallel-hybrid, and pure-Transformer models.
- Prior-art differentiation. We extend DecDEC’s short-horizon Transformer evidence to the long-reasoning regime and show that Quamba2’s channel-persistence assumption fails at the measured block-output surface for the hybrid models tested here.
- Mechanism tests. We evaluate seven channel-set protection families with random and matched-cost controls, yielding a three-mechanism account: boundary discontinuities are harmful, smooth top-1% tracking is insufficient, and protected budget is load-bearing; KL accumulation does not support runaway compound error in the measured Granite packet.
- Partial remedy and scope. Budget-tuned EMA protection succeeds on Nemotron at top-10 and shows partial Granite signal, while ParoQuant recovers 0.75 on Granite via rotation. Thus the negative result is scoped to channel-set protection, not all W4A16 PTQ.

2 Related Work

Mamba and state-space quantization. Vision-Mamba quantization already reports dynamic activation outliers: QMamba introduces temporal grouping for hidden states (Li et al., 2025), and OuroMamba uses dynamic outlier detection during inference (Ramachandran et al., 2025). Language-Mamba work often assumes more stable structure. Mamba-PTQ identifies recurrent outlier channels (Pierro and Abreu, 2024); Quamba and Quamba2 use static or calibration-time structure for Mamba-family PTQ (Chiang et al., 2024; 2025); MambaQuant uses variance-aligned rotations (Xu et al., 2025). Our measurement asks whether the static protected-set assumption holds in long reasoning traces.¹

LLM quantization and rotations. SmoothQuant, AWQ, QuaRot, KVQuant, and BlockDialect protect or transform activation structure using calibration-time statistics, rotations, or adaptive formats (Xiao et al., 2023; Lin et al., 2023; Ashkboos et al., 2024; Hooper et al., 2024; Jang and Tambe, 2025). These systems are not contradicted by our results. The implication is narrower: if the protected channel set changes during long decode, static channel maps need validation on the specific architecture and horizon. ParoQuant is particularly important because it reports strong W4A16 recovery for reasoning models through pairwise rotations (Liang et al., 2025); our ParoQuant baseline contextualizes, rather than replaces, the channel-set experiments.

Dynamic and long-context policies. DecDEC dynamically identifies salient activation channels at each decode step and reports low recall for static outlier-channel identification in short Transformer decode (Park et al., 2025). Our novelty is horizon, model class, architecture, and reasoning benchmark surface. KV-cache methods such as KIVI, PM-KVQ, AttentionPredictor, ChanMix, and MixKVQ adapt or compress cached key/value tensors (Liu et al., 2024; 2025; Yang et al., 2025; Liao and Wen, 2026; Zhang et al., 2025); our interventions target linear-layer activation channels.²

Error accumulation. SLQ finds that near-lossless quantization errors need not compound across token positions (Helcig et al., 2026). QEP, LAQuant, and Activation Sensitivity emphasize layer-wise propagation, lookahead losses, calibration mismatch, and task-conditional sensitivity as PTQ risks (Arai and Ichikawa, 2025; Choi et al., 2026; Xu, 2026). We therefore test, rather than assume, compound error in our aggressive W4A16 endpoint setting.

¹We use decode-time channel-set drift for the measured phenomenon. SmoothQuant’s “migration” is activation-to-weight scale transfer (Xiao et al., 2023); MoBiQuant uses related language for token-set shifts across bit-widths (Wang et al., 2026).

²“Hot channels” in HCP/CHON refers to persistently extreme channels during NVFP4 pretraining, which differs from our W4A16 inference-time setting. We refer to arXiv:2603.19296 as test-time TTQ to avoid conflict with classical Trained Ternary Quantization.

3 Method

For each trace, layer, and measured decode position, we record the vector of per-channel activation magnitudes. Let $S_{\ell,p}(t)$ be the top-1% channel set for layer ℓ and prompt p at position t . The strict set-leaving rate at position t is

$$1 - \frac{|S_{\ell,p}(100) \cap S_{\ell,p}(t)|}{|S_{\ell,p}(100)|}, \quad (1)$$

averaged over prompts and layers. Within-set shuffling is reported separately for channels that remain in the top-1% set but move substantially in rank.

We score interventions using recovery of the BF16-vs-static W4A16 gap:

$$R = \frac{L_{\text{static}} - L_{\text{method}}}{L_{\text{static}} - L_{\text{BF16}}}. \quad (2)$$

Recovery may be negative when a method is worse than static protection and may exceed one when a method beats BF16 on the scored slice. We therefore report unclipped medians and bootstrap CIs. Traces with no recoverable static gap are tracked explicitly and excluded from ratio computation for method CIs; their frequency is part of the trace-heterogeneity result.

We test seven channel-set method classes: static migration-aware unions, hard position-conditioned protected sets, hard position-binned scale tables, EMA-smoothed top-magnitude protection (M11), cross-tensor activation+key coupling (M18), DecDEC-style reactive top-1% selection, budget-tuned EMA protection (M11b), and stable-core protection (M26). Each method is compared to static baselines plus random or matched-budget controls when applicable. ParoQuant is evaluated as an algorithmic rotation baseline, not as a channel-set method.

4 Experiments

Models and traces. The measurement suite covers Granite-4.0-H-Small, Nemotron-3-Nano, DeepSeek-R1-Distill-Qwen-1.5B, and Falcon-H1. Intervention experiments use 12 deterministic long-reasoning traces unless otherwise noted; the original drift measurement uses 24 Granite-Small traces. Decode positions are measured on the grid $\{100, 500, 1000, 5000, 10000, 20000\}$ when the model packet supports 20K tokens. Dense KL analysis uses a Granite-Small packet sampled every 100 positions to 10K.

Controls. Static top-1% protection is the baseline. Position- and scale-switching methods use random-bin controls. EMA and reactive methods use random-walk or random-reactive controls. Budget experiments include matched static top-10 controls. This control design is load-bearing: two methods that look plausible in isolation are worse than their random controls.

5 Results

5.1 Channel sets drift at long reasoning horizons

Figure 1 shows the headline measurement. All four models exhibit increasing strict set-leaving across decode position. By the final measured position, strict set-leaving is 0.566 for Granite-Small, 0.534 for Nemotron-3, 0.671 for DeepSeek-R1-Distill, and 0.674 for Falcon-H1. The pure-Transformer reference drifts at least as much as the hybrid models, so the phenomenon is not confined to Mamba layers.

5.2 Layer type does not explain the drift away

The Quamba2 contradiction is strongest if drift appears in SSM/Mamba block outputs, not only in attention or MLP outputs. Figure 2 shows that Granite and Nemotron block-output




Figure 1: Strict top-1% channel-set leaving relative to decode position 100. Static channel protection loses 53–67% of its original protected set by the final measured position across four reasoning models.




Figure 2: Layer-stratified drift at measured block outputs. Strict set-leaving is comparable across attention, SSM/Mamba, and MoE-classified layers, so the drift is not localized to one component type.

drift is comparable across layer types. Granite attention and SSM/Mamba outputs have strict set-leaving of 0.557 and 0.567. Nemotron attention, MoE, and SSM/Mamba outputs have 0.563, 0.527, and 0.533. This does not isolate internal SSM state tensors, but it does show that the measured block-output surface is not channel-persistent.

5.3 Channel-set interventions mostly fail at top-1% budget

Figure 3 and Table 1 summarize the intervention outcomes. Hard position-conditioned protection is outperformed by its random-bin control by 0.668 median recovery, and hard scale bins are outperformed by random scale-bin assignment by 0.761. EMA smoothing is different: M11 beats its random-walk control by 1.25 median, but its absolute recovery is only 0.048. M18 activation+key coupling has the strongest signal-vs-random margin among failed top-1% methods, beating random coupling by 4.32 median, but its absolute recovery is negative. DecDEC-style reactive selection is less bad than a random reactive control




Figure 3: Median recovery with 95% bootstrap CIs. Long negative CI whiskers are clipped at the plot boundary for readability; exact rounded values are in Table 1. The dashed line marks the 0.30 positive-method threshold used in preregistered gates.

Method family	Model	Median recovery	95% CI	Key comparison
Position-conditioned channels	Granite	-0.867	[-3.44, 0.595]	random bin better by 0.668
Position-binned scales	Granite	0.234	[-2.07, 0.504]	random scale bin better by 0.761
EMA-smoothed top-1%	Granite	0.048	[-8.94, 0.705]	random-walk lower by 1.25
Activation+key coupling	Granite	-0.344	[-14.1, 0.741]	random coupled lower by 4.32
DecDEC proxy	Granite	-0.070	[-8.50, 0.674]	random reactive lower by 1.19
Stable core	Granite	0.178	[-0.116, 1.00]	random matched-core lower by 1.73
Budget EMA top-5	Granite	0.449	[-1.30, 1.00]	static top-10 lower by 0.501
Budget EMA top-5	Nemotron	0.457	[0.346, 0.795]	static top-10 higher by 0.138
Budget EMA top-10	Nemotron	0.815	[0.254, 0.923]	static top-10 lower by 0.220
ParoQuant rotation	Granite	0.754	[0.477, 1.00]	M11b top-5 lower by 0.304

Table 1: Intervention outcomes. Recovery is unclipped; values can be negative or exceed one. ParoQuant is included as a rotation baseline and is not a channel-set method.

but still does not recover quality. M26 stable-core protection shows signal over random matched-core channels but remains below the positive-method threshold.

5.4 Budget is load-bearing, but not model-invariant

M11b changes the conclusion from uniformly negative to partially positive. On Granite-Small, top-5 EMA protection recovers 0.449 of the static gap and beats static top-10 by 0.501 median, but the CI is wide. On Nemotron, top-5 is positive with a strictly positive CI, yet it does not beat static top-10. The top-10 arm is the strongest channel-set result: median recovery 0.815, CI [0.254, 0.923], and a 0.220 median margin over static top-10. Thus budget is a real mechanism, but the best budget and the value over static protection are model-dependent.

ParoQuant constrains the claim. Its Granite-Small rotation baseline recovers 0.754, exceeding M11b top-5 on the same model. The paper therefore does not argue that W4A16 PTQ is intrinsically broken; it argues that channel-set protection is fragile under long reasoning drift, while rotation-based conditioning can address part of the same endpoint by a different mechanism.




Figure 4: Per-position KL trajectories on Granite-Small. All three regimes prefer sublinear square-root fits and have moderate AR decay, rather than superlinear growth.

5.5 KL does not support runaway compounding

Figure 4 tests whether failed channel-set methods primarily fail because errors compound over decode. In the dense Granite-Small packet, static top-1%, DecDEC proxy, and M11 all select a sublinear square-root fit over linear or superlinear alternatives. Mean KL is 0.150, 0.133, and 0.131 respectively, and fitted AR decays are 0.51, 0.45, and 0.44. This weakens compound error as the primary explanation in this measured regime and shifts attention back to stale or misbudgeted protection.

6 Discussion

The results support a three-mechanism account. First, boundary discontinuities are actively harmful: both hard protected-set and hard scale-bin switching lose to random controls. Second, smoothing removes the worst boundary harm but does not recover quality at top-1% budget; M11 and M18 contain signal, but not enough protected capacity. Third, budget is binding, but budget alone is not a universal recipe: Nemotron top-10 succeeds, Granite top-5 is promising but high-variance, and static top-10 is already strong on Nemotron.

Trace heterogeneity remains important. In the Granite channel-set experiments, roughly one third of the 12-trace slice has no recoverable static gap under the ratio metric; in the larger static-union experiment the no-gap fraction is 0.375. This limits median-based claims and explains the wide CIs on several methods. The paper therefore emphasizes controlled mechanism comparisons and cross-model directionality rather than a single universal method headline.

The natural next step is composition rather than another threshold-only channel set. ParoQuant shows that rotations can recover much of the Granite W4A16 gap; M11b shows that protected budget matters. A stronger ICLR follow-up would test rotation plus budget-aware protection, larger reasoning scales, and training-time or calibration-time interventions that make long-decode channel statistics less brittle.

7 Limitations

The scope is W4A16 PTQ on small reasoning models and measured block-output activation channels. We do not claim results for FP8, MXFP4/NVFP4, W4A4, or quantization-aware training. The hybrid-model Quamba2 contradiction is at the accessible block-output surface, not at every internal SSM state tensor. Only Granite-Small has dense FFT and

KL diagnostics; four-model dense spectral and per-position KL replication remain future work. The M11b positive result is cross-model but not model-invariant: Nemotron top-10 succeeds, while Granite’s strongest signal is top-5 with wide uncertainty.

8 Conclusion

Long-reasoning W4A16 channel protection fails when static protected sets drift, and our experiments show that this drift is substantial across four measured models. Budget-tuned smoothing partially remedies the failure, especially on Nemotron, but rotations and broader calibration changes remain the stronger path toward a deployable method.

References

- Yamato Arai and Yuma Ichikawa. Quantization Error Propagation: Revisiting Layer-Wise Post-Training Quantization. arXiv:2504.09629, NeurIPS 2025. <https://arxiv.org/abs/2504.09629>.
- Julie Bort. The rise of AI reasoning models is making benchmarking more expensive. TechCrunch, April 10 2025. <https://techcrunch.com/2025/04/10/the-rise-of-ai-reasoning-models-is-making-benchmarking-more-expensive/>.
- Gartner. Gartner predicts that by 2030, performing inference on an LLM with 1 trillion parameters will cost GenAI providers over 90% less than in 2025. Press release, March 25 2026. <https://www.gartner.com/en/newsroom/press-releases/2026-03-25-gartner-predicts-that-by-2030-performing-inference-on-an-llm-with-1-trillion-parameters-will-cost-genai-providers-over-90-percent-less-than-in-2025>.
- MarketsandMarkets. AI inference market size, share and growth, 2025 to 2030. Commercial market report, February 2025. <https://www.marketsandmarkets.com/Market-Reports/ai-inference-market-189921964.html>.
- MarketsandMarkets. Edge AI hardware industry worth \$58.90 billion by 2030. Commercial market report press release, July 4 2025. <https://www.marketsandmarkets.com/PressReleases/edge-ai-hardware.asp>.
- Dongwei Wang, Jinhee Kim, Seokho Han, Denis Gudovskiy, Yohei Nakata, Tomoyuki Okuno, KhayTze Peong, Kang Eun Jeon, Jong Hwan Ko, Yiran Chen, and Huanrui Yang. MoBiQuant: Mixture-of-Bits Quantization for Token-Adaptive Elastic LLMs. arXiv:2602.20191, February 2026. <https://arxiv.org/abs/2602.20191>.
- Yeonhong Park, Jake Hyun, Hojoon Kim, and Jae W. Lee. DecDEC: A Systems Approach to Advancing Low-Bit LLM Quantization. 19th USENIX Symposium on Operating Systems Design and Implementation (OSDI 25), 2025. <https://www.usenix.org/system/files/osdi25-park-yeonhong.pdf>.
- Zirui Liu, Jiayi Yuan, Hongye Jin, Shaochen Zhong, Zhaozhao Xu, Vladimir Braverman, Beidi Chen, and Xia Hu. KIVI: A Tuning-Free Asymmetric 2bit Quantization for KV Cache. ICML 2024. <https://arxiv.org/abs/2402.02750>.
- Tengxuan Liu, Shiyao Li, Jiayi Yang, Tianchen Zhao, Feng Zhou, Xiaohui Song, Guohao Dai, Shengen Yan, Huazhong Yang, and Yu Wang. PM-KVQ: Progressive Mixed-precision KV Cache Quantization for Long-CoT LLMs. arXiv:2505.18610, 2025. <https://arxiv.org/abs/2505.18610>.
- Qingyue Yang, Jie Wang, Xing Li, Zhihai Wang, Chen Chen, Lei Chen, Xianzhi Yu, Wulong Liu, Jianye Hao, Mingxuan Yuan, and Bin Li. AttentionPredictor: Temporal Patterns Matter for KV Cache Compression. NeurIPS 2025. <https://arxiv.org/abs/2502.04077>.
- Hao Mark Chen, Fuwen Tan, Alexandros Kouris, Royson Lee, Hongxiang Fan, and Stylianos I. Venieris. Progressive Mixed-Precision Decoding for Efficient LLM Inference. arXiv:2410.13461, revised February 2025. <https://arxiv.org/abs/2410.13461>.

-
- Yinglong Li, Xiaoyu Liu, Jiacheng Li, Ruikang Xu, Yinda Chen, and Zhiwei Xiong. QMamba: Post-Training Quantization for Vision State Space Models. arXiv:2501.13624, January 2025. <https://arxiv.org/abs/2501.13624>.
- Akshat Ramachandran, Mingyu Lee, Huan Xu, Souvik Kundu, and Tushar Krishna. Ouro-Mamba: A Data-Free Quantization Framework for Vision Mamba. arXiv:2503.10959, ICCV 2025. <https://arxiv.org/abs/2503.10959>.
- Alessandro Pierro and Steven Abreu. Mamba-PTQ: Outlier Channels in Recurrent Large Language Models. arXiv:2407.12397, 2024. <https://arxiv.org/abs/2407.12397>.
- Hung-Yueh Chiang, Chi-Chih Chang, Natalia Frumkin, Kai-Chiang Wu, and Diana Marculescu. Quamba: A Post-Training Quantization Recipe for Selective State Space Models. arXiv:2410.13229, 2024. <https://arxiv.org/abs/2410.13229>.
- Hung-Yueh Chiang, Chi-Chih Chang, Natalia Frumkin, Kai-Chiang Wu, Mohamed S. Abdelfattah, and Diana Marculescu. Quamba2: A Robust and Scalable Post-training Quantization Framework for Selective State Space Models. arXiv:2503.22879, 2025. <https://arxiv.org/abs/2503.22879>.
- Euntae Choi, Sumin Song, and Sungjoo Yoo. LAQuant: A Simple Overhead-free Large Reasoning Model Quantization by Layer-wise Lookahead Loss. arXiv:2605.08755, 2026. <https://arxiv.org/abs/2605.08755>.
- Zukang Xu, Yuxuan Yue, Xing Hu, Zhihang Yuan, Zixu Jiang, Zhixuan Chen, Jiangyong Yu, Chen Xu, Sifan Zhou, and Dawei Yang. MambaQuant: Quantizing the Mamba Family with Variance Aligned Rotation Methods. arXiv:2501.13484, 2025. <https://arxiv.org/abs/2501.13484>.
- Yizhi Chen and Ahmed Hemani. Late Breaking Results: Quamba-SE: Soft-edge Quantizer for Activations in State Space Models. arXiv:2601.09451, DATE Late Breaking Results 2026. <https://arxiv.org/abs/2601.09451>.
- Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models. arXiv:2211.10438, ICML 2023. <https://arxiv.org/abs/2211.10438>.
- Ji Lin, Jiaming Tang, Haotian Tang, Shang Yang, Wei-Ming Chen, Wei-Chen Wang, Guangxuan Xiao, Xingyu Dang, Chuang Gan, and Song Han. AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration. arXiv:2306.00978, MLSys 2024. <https://arxiv.org/abs/2306.00978>.
- Saleh Ashkboos, Amirkeivan Mohtashami, Maximilian L. Croci, Bo Li, Pashmina Cameron, Martin Jaggi, Dan Alistarh, Torsten Hoefer, and James Hensman. QuaRot: Outlier-Free 4-Bit Inference in Rotated LLMs. arXiv:2404.00456, 2024. <https://arxiv.org/abs/2404.00456>.
- Coleman Hooper, Sehoon Kim, Hiva Mohammadzadeh, Michael W. Mahoney, Yakun Sophia Shao, Kurt Keutzer, and Amir Gholami. KVQuant: Towards 10 Million Context Length LLM Inference with KV Cache Quantization. arXiv:2401.18079, NeurIPS 2024. <https://arxiv.org/abs/2401.18079>.
- Wonsuk Jang and Thierry Tambe. BlockDialect: Block-wise Fine-grained Mixed Format Quantization for Energy-Efficient LLM Inference. arXiv:2501.01144, ICML 2025. <https://arxiv.org/abs/2501.01144>.
- Michael Helcig, Eldar Kurtic, and Dan Alistarh. Statistically-Lossless Quantization of Large Language Models. arXiv:2605.02404, 2026. <https://arxiv.org/abs/2605.02404>.
- Yesheng Liang, Haisheng Chen, Song Han, and Zhijian Liu. ParoQuant: Pairwise Rotation Quantization for Efficient Reasoning LLM Inference. arXiv:2511.10645, 2025. <https://arxiv.org/abs/2511.10645>.

-
- Chengxi Liao and Zeyi Wen. Channel-Aware Mixed-Precision Quantization for Efficient Long-Context Inference. ICLR 2026. <https://openreview.net/forum?id=yjr2jX41qO>.
- Qwen Team. Qwen3.6-35B-A3B model card. Hugging Face, 2026. <https://huggingface.co/Qwen/Qwen3.6-35B-A3B>.
- Bruce Changlong Xu. Activation Sensitivity as a Unifying Principle for Post-Training Quantization. arXiv:2601.11663, 2026. <https://arxiv.org/abs/2601.11663>.
- Songlin Yang, Bailin Wang, Yikang Shen, Rameswar Panda, and Yoon Kim. Gated Linear Attention Transformers with Hardware-Efficient Training. arXiv:2312.06635, 2023. <https://arxiv.org/abs/2312.06635>.
- Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated Delta Networks: Improving Mamba2 with Delta Rule. arXiv:2412.06464, 2024. <https://arxiv.org/abs/2412.06464>.
- Moonshot AI Kimi Team: Yu Zhang et al. Kimi Linear: An Expressive, Efficient Attention Architecture. arXiv:2510.26692, October 2025. <https://arxiv.org/abs/2510.26692>.
- Tao Zhang, Ziqian Zeng, Hao Peng, Huiping Zhuang, and Cen Chen. MixKVQ: Query-Aware Mixed-Precision KV Cache Quantization for Long-Context Reasoning. arXiv:2512.19206, 2025. <https://arxiv.org/abs/2512.19206>.

Ethics Statement

This work analyzes benchmark prompts and model activations. It does not use private user data or make deployment-safety claims. The main ethical risk is overclaiming a scoped mechanism study as a universal quantization recipe; the paper therefore separates measurement, mechanism, and partial-remedy claims.

Reproducibility Statement

The release package contains the public reproduction interface, including `release/README.md`, `release/docs/reproducing_results.md`, and `release/scripts/reproduce_all.sh`. Full-fidelity reproduction reruns every paper-referenced method and analysis on a single high-memory GPU; fast verification mode checks the same code paths with reduced trace counts.

Model	HuggingFace identifier	Checkpoint note
Granite-Small	<code>ibm-granite/granite-4.0-h-small</code>	pinned release snapshot
Nemotron-3	<code>nvidia/NVIDIA-Nemotron-3-Nano-30B-A3B-BF16</code>	pinned release snapshot
DeepSeek-R1-Distill	<code>deepseek-ai/DeepSeek-R1-Distill-Qwen-1.5B</code>	pinned release snapshot
Falcon-H1	<code>tiuae/Falcon-H1-0.5B-Instruct</code>	pinned release snapshot

Table 2: Model checkpoints used by the release package. Exact snapshot commits and file hashes are recorded in the release documentation to keep the main paper readable.

Bootstrap seeds, deterministic trace indices, method budgets, and tolerance windows are specified in the release configs. GPU kernels can still exhibit minor nondeterminism across driver versions; the verification compares outputs to bootstrap CI tolerances rather than bit-exact scores. The original Qwen3.6 exploratory hook surface was not used for paper claims because activation access was incomplete.

A Detailed Experimental Provenance

The valid Nemotron budget-tuning result uses the corrected static top-1% baseline rerun and reuses validated score caches for the budget arms. The earlier invalid baseline is excluded from all recovery claims. The release mapping table records both the valid rerun and the excluded failed attempt so that reviewers can audit the correction.

Paper description	Internal study label	Provenance location
Granite-Tiny pilot drift measurement	Phase 0	experimental/outlier_migrate/phase0/
Granite-Small drift replication	Phase 1	experimental/outlier_migrate/phase1/
Nemotron-3 drift replication	Phase 2	experimental/outlier_migrate/phase2/
Static migration-aware union, Granite-Tiny	Phase 3	experimental/outlier_migrate/phase3/
Static migration-aware union, Granite-Small	Phase 4	experimental/outlier_migrate/phase4/
DeepSeek Transformer control	Phase 5 prime	experimental/outlier_migrate/phase5_prime/
Falcon-H1 parallel-hybrid control	Phase 7	experimental/outlier_migrate/phase7/
Channel-set methods and baselines	Method study	experimental/outlier_migrate/phase9/
Paper release interface	Release package	release/docs/reproducing_results.md

Table 3: Translation from paper-facing experiment names to internal study labels. Exact packet directories, command lines, prompt hashes, and model file hashes are maintained in the release documentation and result manifests.

B Per-Method Detailed Results

Experiment	Included traces	No-gap fraction	Median recovery	95% CI
Static union, Granite-Tiny	24	0.417	0.000	[0.000, 0.711]
Static union, Granite-Small	24	0.375	0.000	[0.000, 0.0695]
Position-conditioned channels	8	0.333	-0.867	[-3.44, 0.595]
Position-binned scales	12	0.000	0.234	[-2.07, 0.504]
EMA-smoothed top-1%	8	0.333	0.048	[-8.94, 0.705]
Activation+key coupling	8	0.333	-0.344	[-14.1, 0.741]
DecDEC proxy	8	0.333	-0.070	[-8.50, 0.674]
Budget EMA top-5, Granite	8	0.333	0.449	[-1.30, 1.00]
Budget EMA top-10, Nemotron	10	0.167	0.815	[0.254, 0.923]
Stable-core protection	8	0.333	0.178	[-0.116, 1.00]
ParoQuant rotation	8	0.333	0.754	[0.477, 1.00]

Table 4: Detailed rounded outcomes for all method families discussed in the main text. Values are recovery ratios relative to each experiment’s BF16-vs-static baseline.

C Extended Layer-Stratified Results

Model	Layer type	Layers	Strict set-leaving	Within-set shuffling
Granite-Tiny	Attention	4	0.618	0.186
Granite-Tiny	SSM/Mamba	36	0.636	0.174
Granite-Small	Attention	4	0.557	0.282
Granite-Small	SSM/Mamba	36	0.567	0.270
Nemotron-3	Attention	6	0.563	0.255
Nemotron-3	MoE	23	0.527	0.278
Nemotron-3	SSM/Mamba	23	0.533	0.264

Table 5: Layer-type decomposition. Strict set-leaving is the component most directly relevant to protected-set membership; within-set shuffling motivates scale-refresh mechanisms but is not itself membership loss.